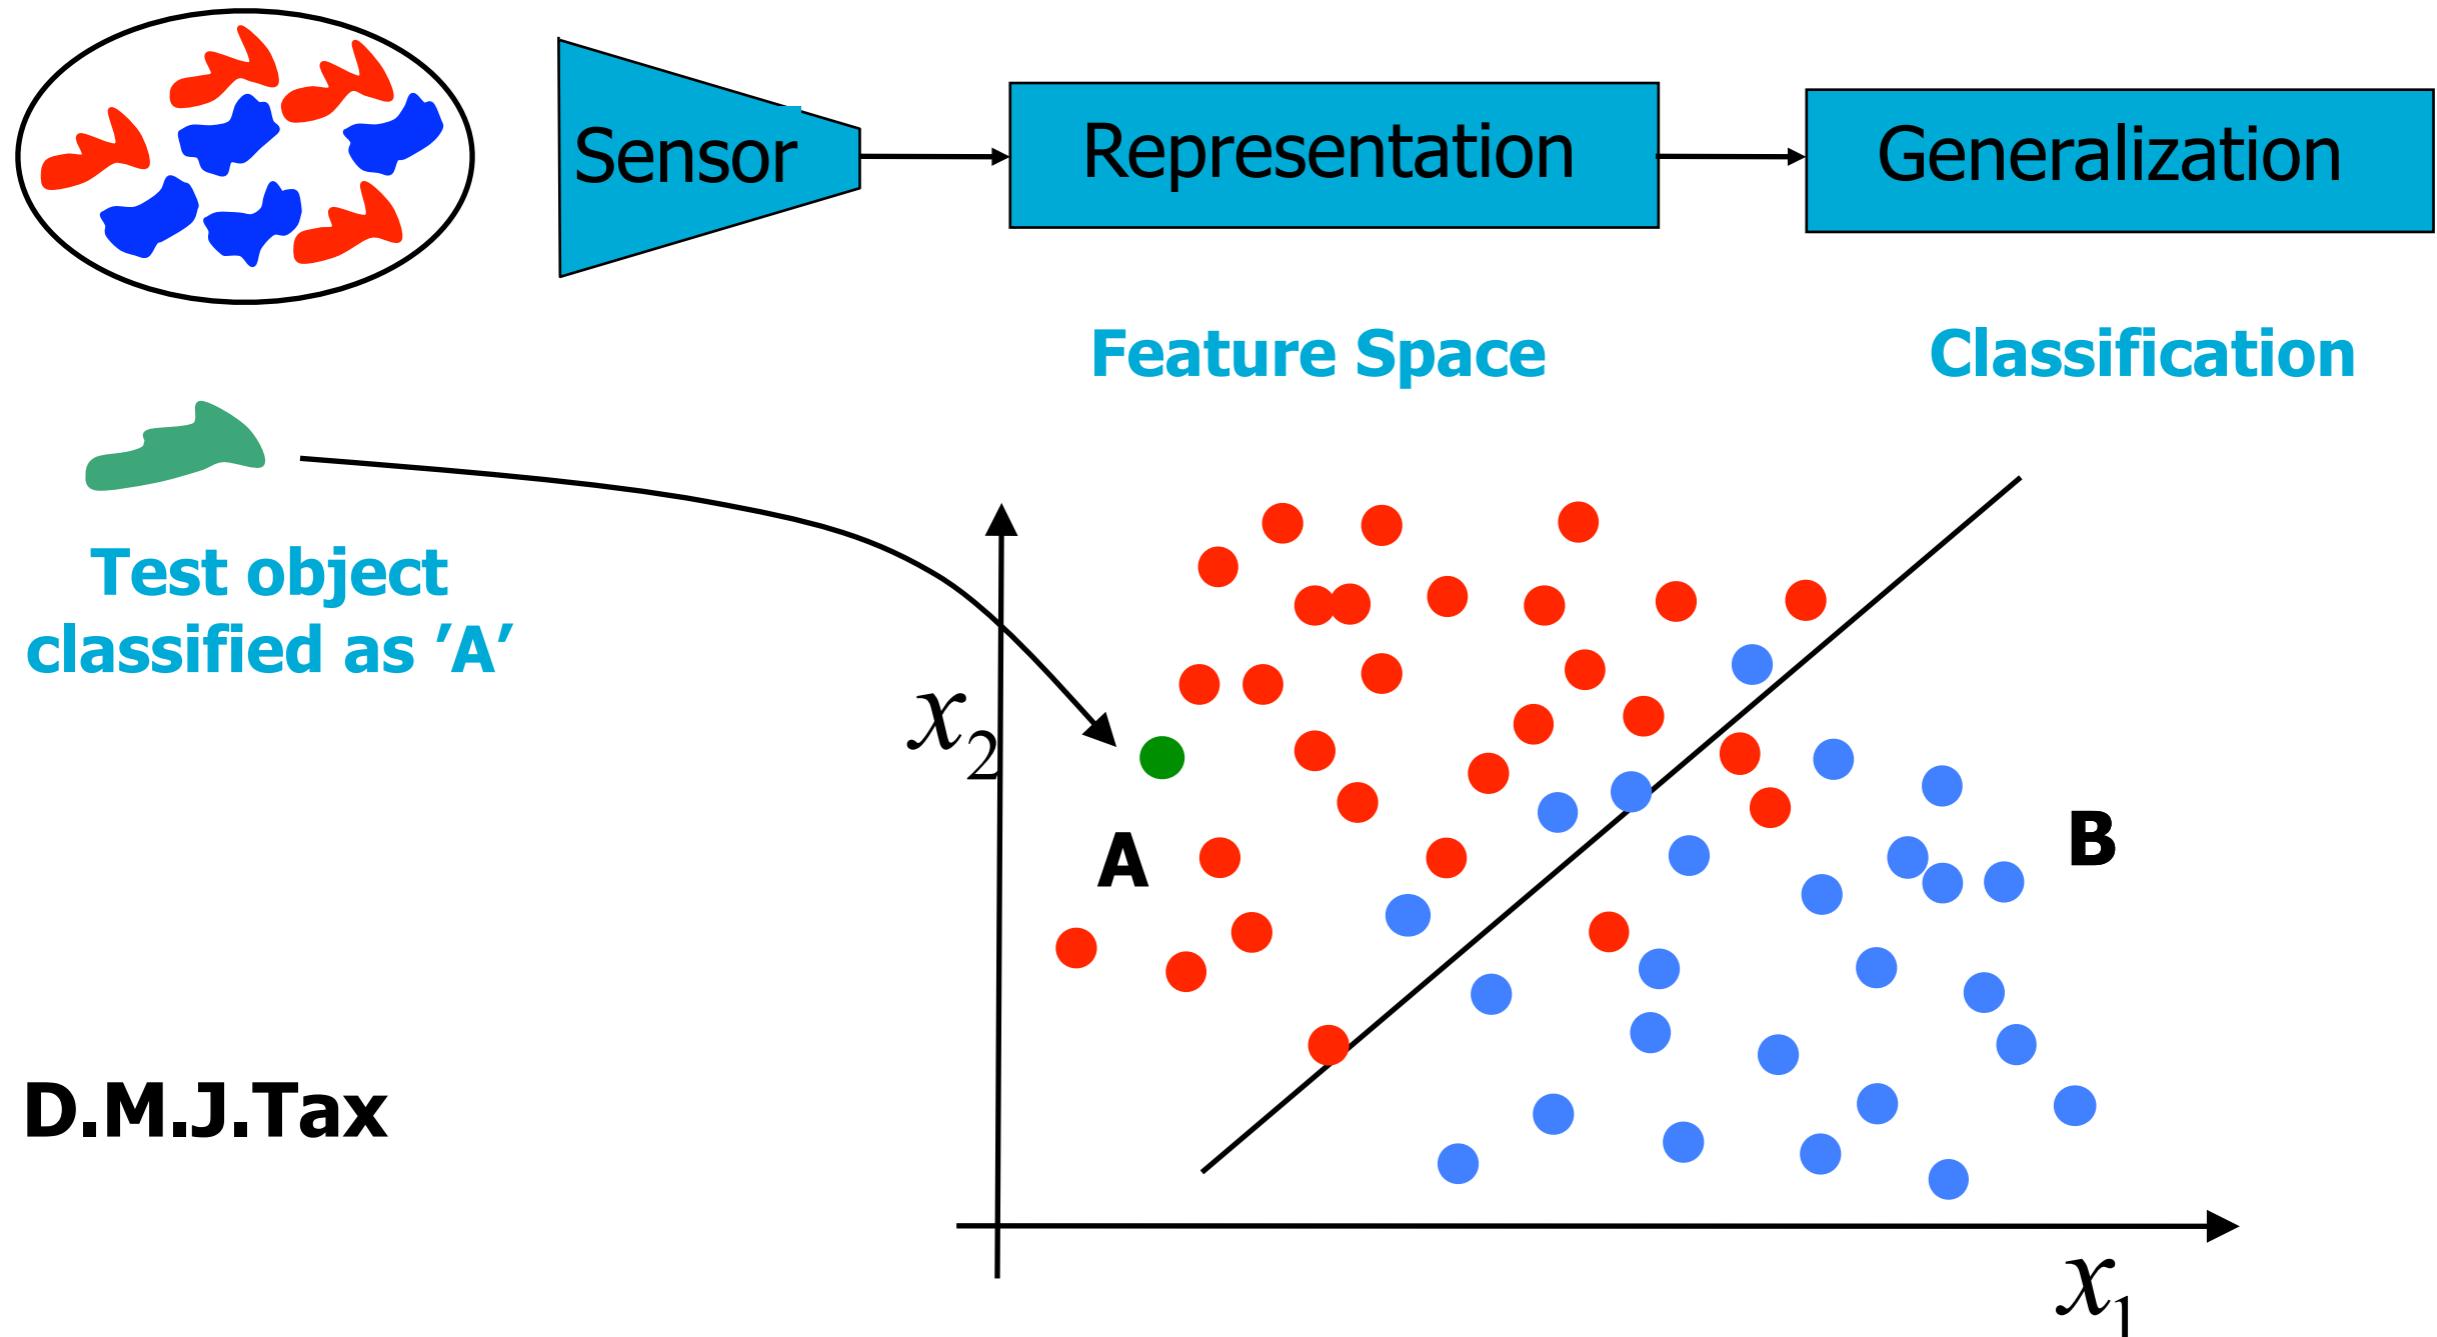


Machine Learning Experiments



Contents

- Hour 1: Some small things (use this for your project as well!):
 - Hints for practical data analysis
 - Stratified cross-validation, leave-one-group-out
 - Paired t-test
 - Hyperparameter optimisation
 - Presentation of results
- Hour 2: Fairness in ML
 - Intuition
 - Formalization of fairness criteria
 - Issues

The final project

- Submit the report and code to Brightspace
- The deadline is mentioned in the Assignments (January 30, 2026)
- Try to make it a Machine Learning problem, not a data analysis problem
- Have a look at the chapter 'Final Project' in the exercise manual
- Focus on the comparison between ML methods, their strengths and weaknesses, and not so much on solving one particular classification/regression problem

Trial Exam

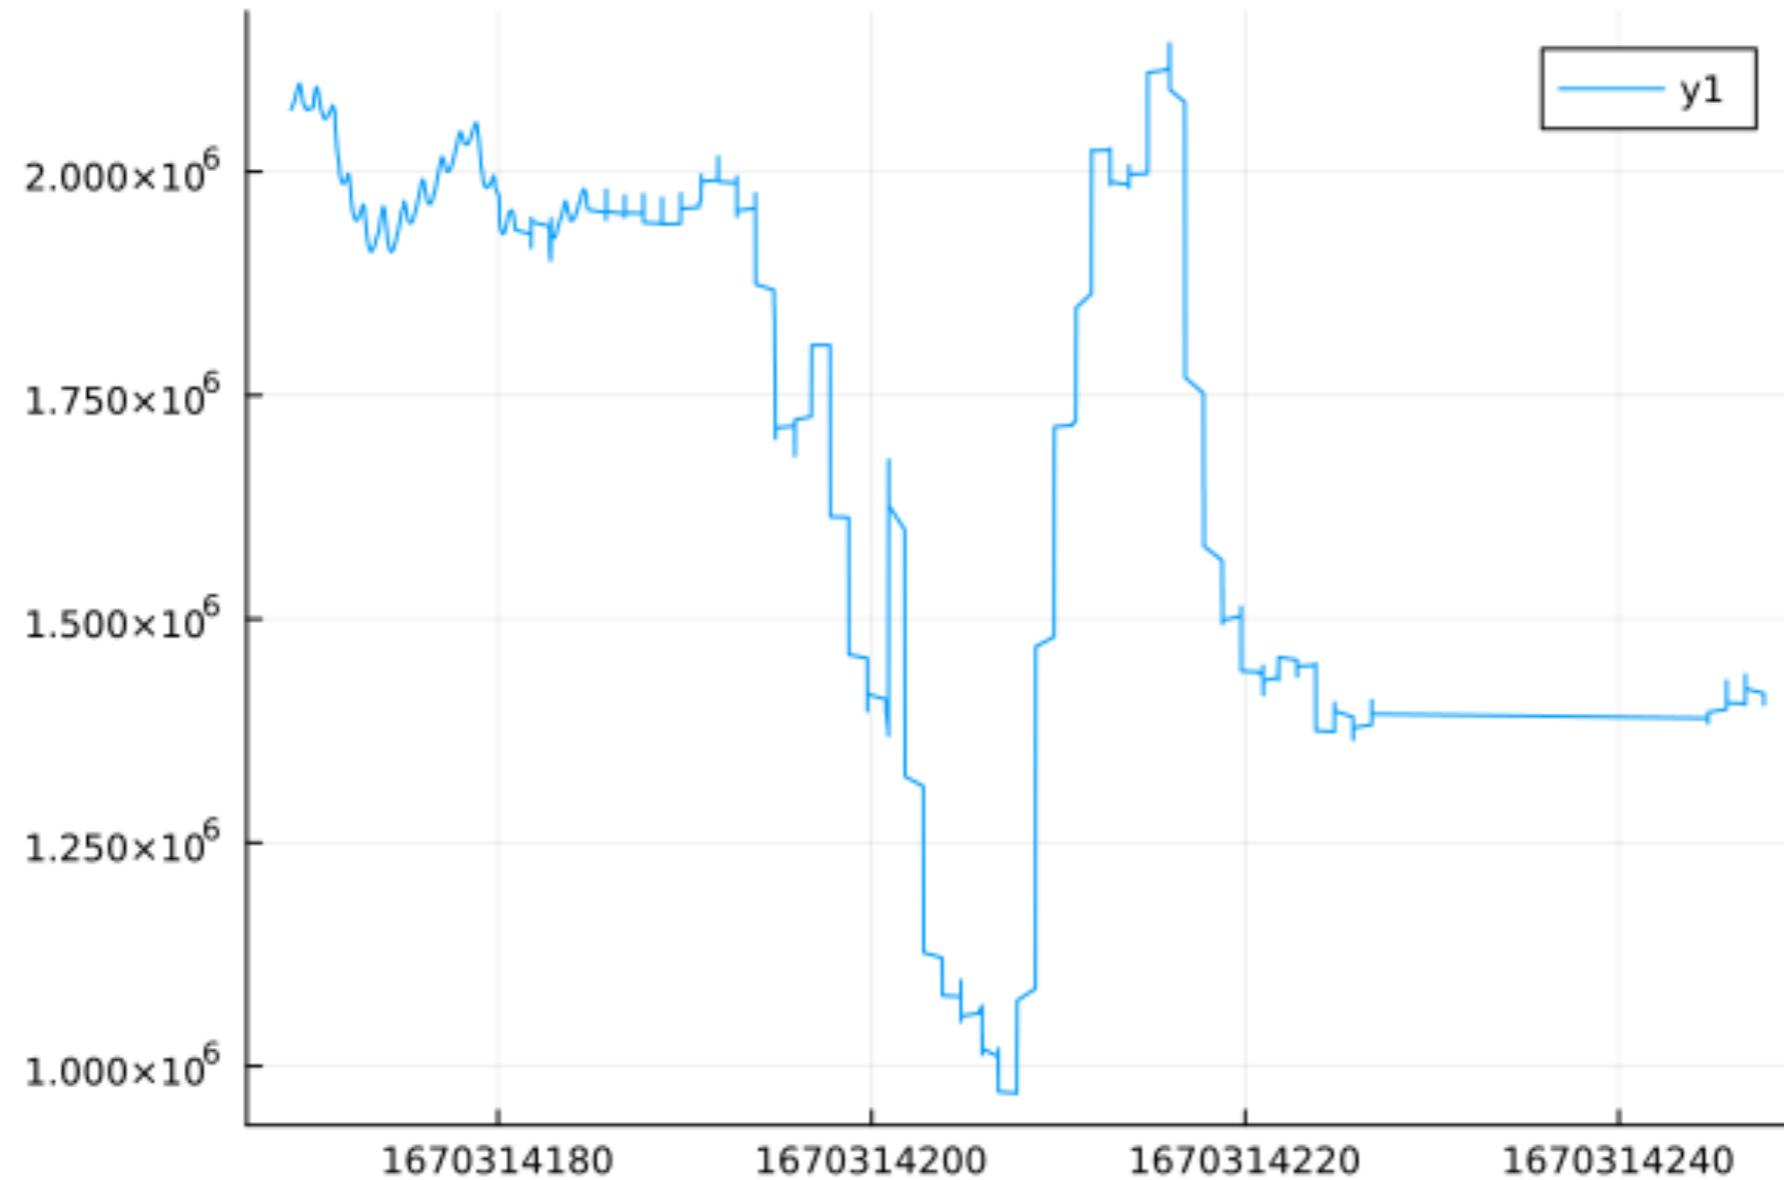
- On weblab.tudelft.nl a practice exam is available
- Try it!
- In the lecture in January we'll have a Q&A session, we can also take these trial exams

Practical data analysis

- Define what you want to predict, the target
- Split your data in a training/design set, and a test set (for smaller datasets, use crossvalidation)
- Eyeball your data (outliers, weird stuff), visualise with scatterplots
- Scale features, invent new features
- Try simple models (linear, nearest neighbour), if needed, extend to more complex
- Be paranoid, use sanity checks
- Which objects are wrongly predicted? Which classes?
- Do a sensitivity analysis ("ablation study")

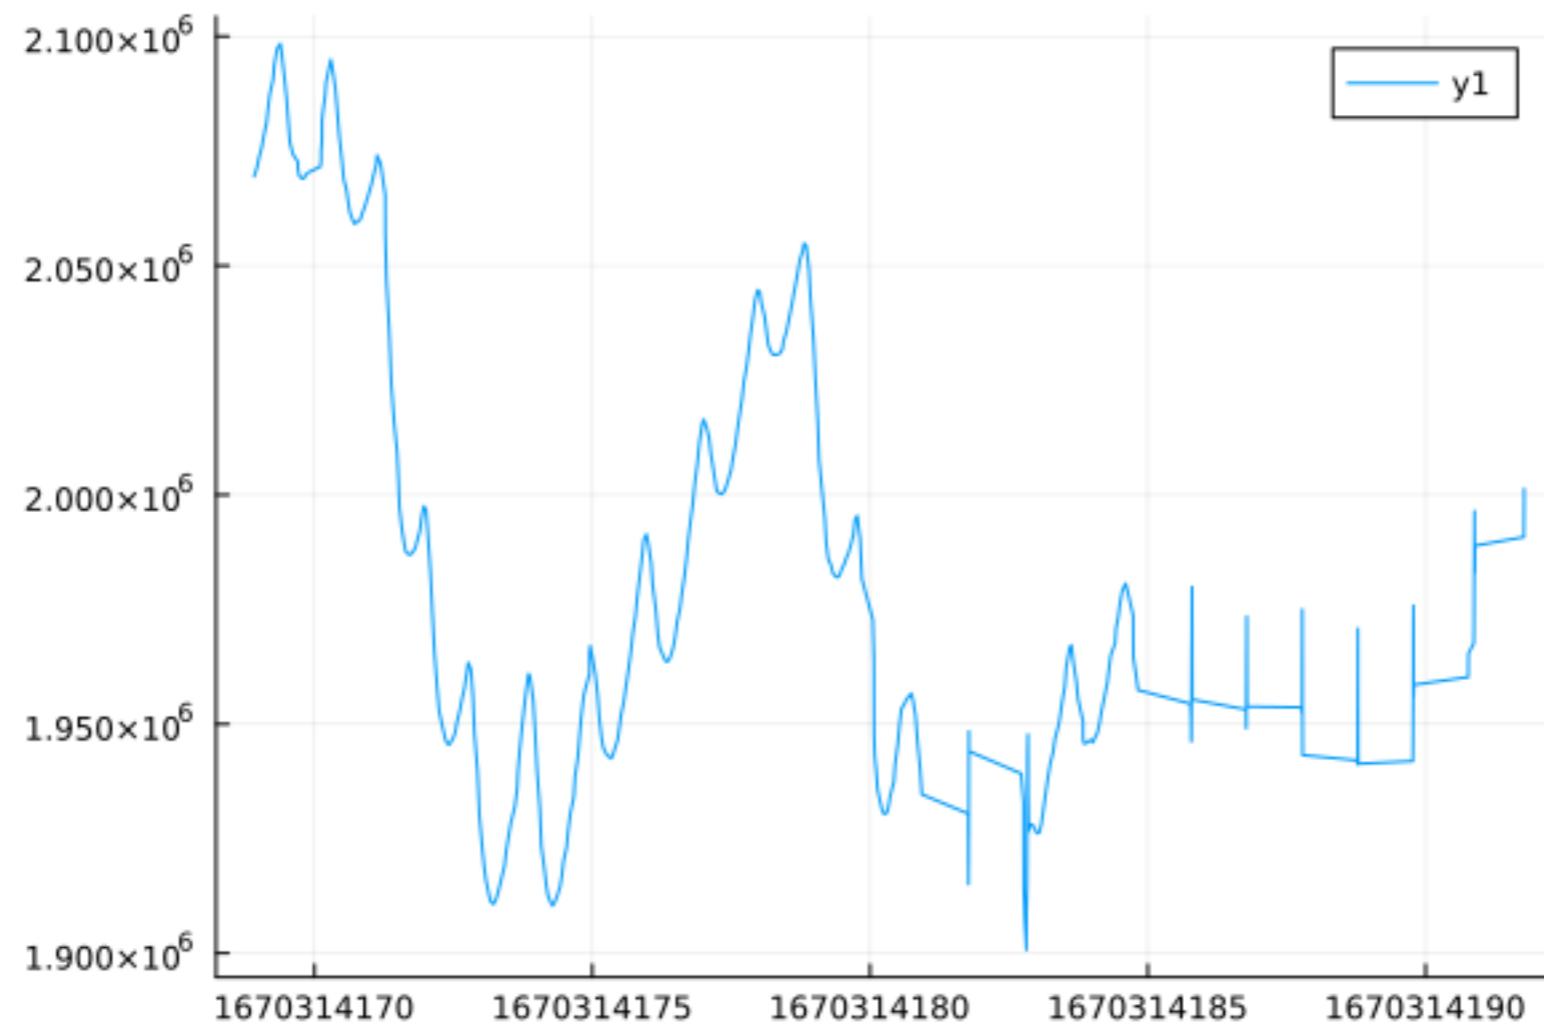
Inspect your data

- Data from Fitbit wearable; a few minutes of PPG signal:



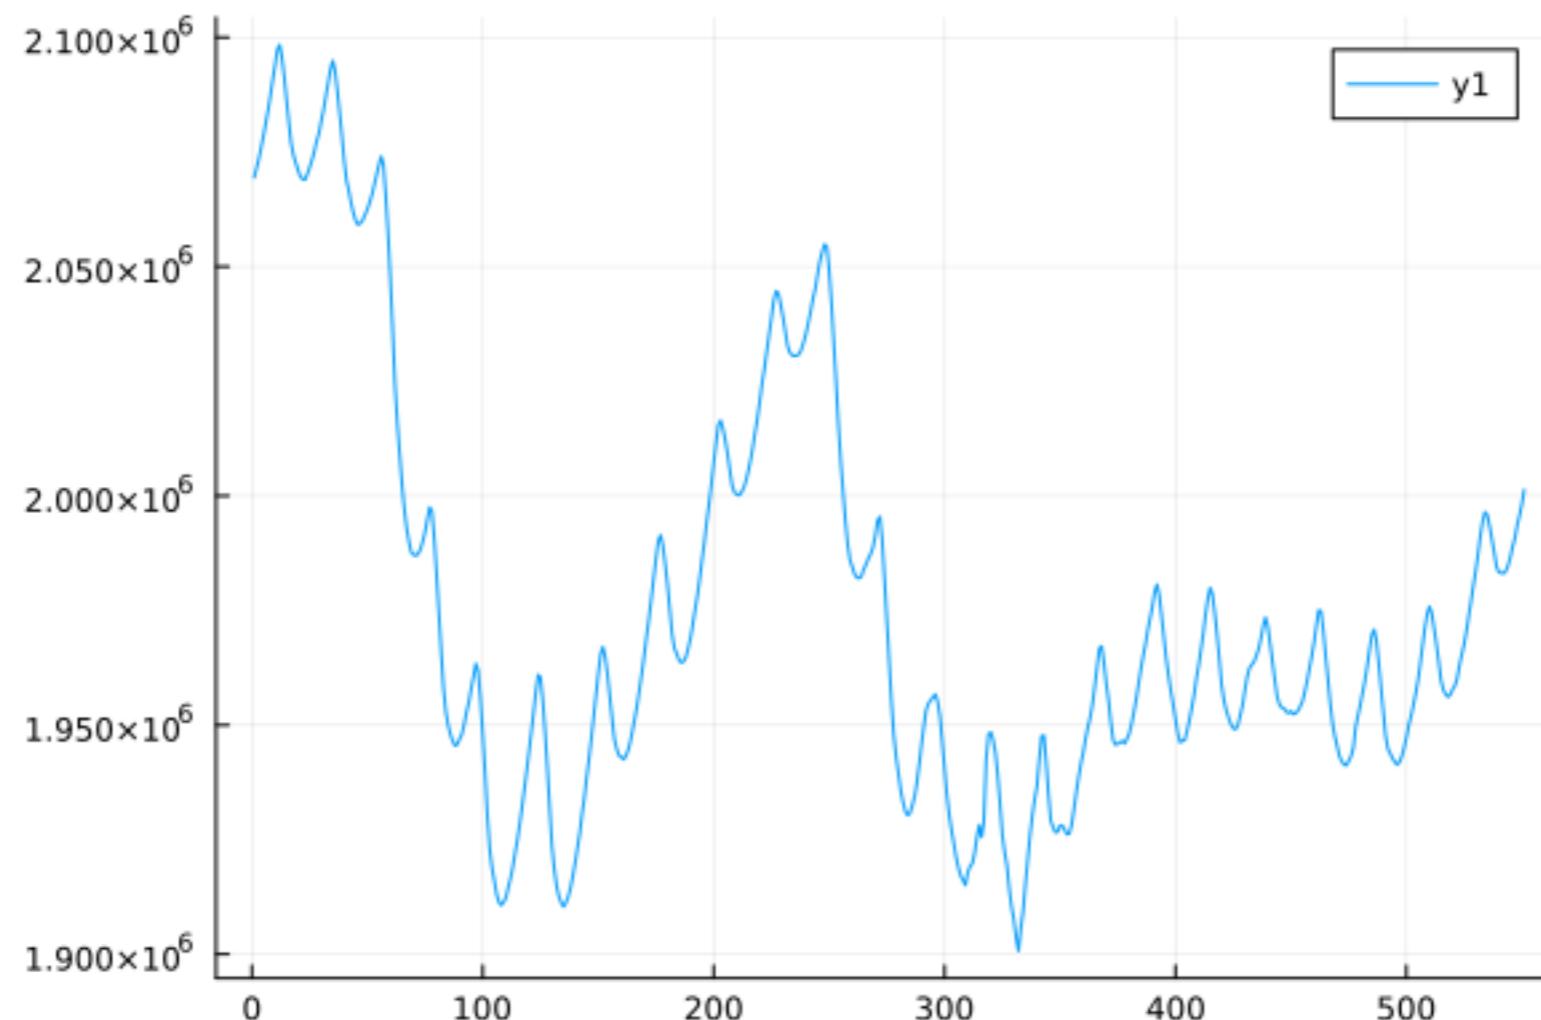
Zoom in of Fitbit data

- First part looks ok-ish, but after 10 seconds...



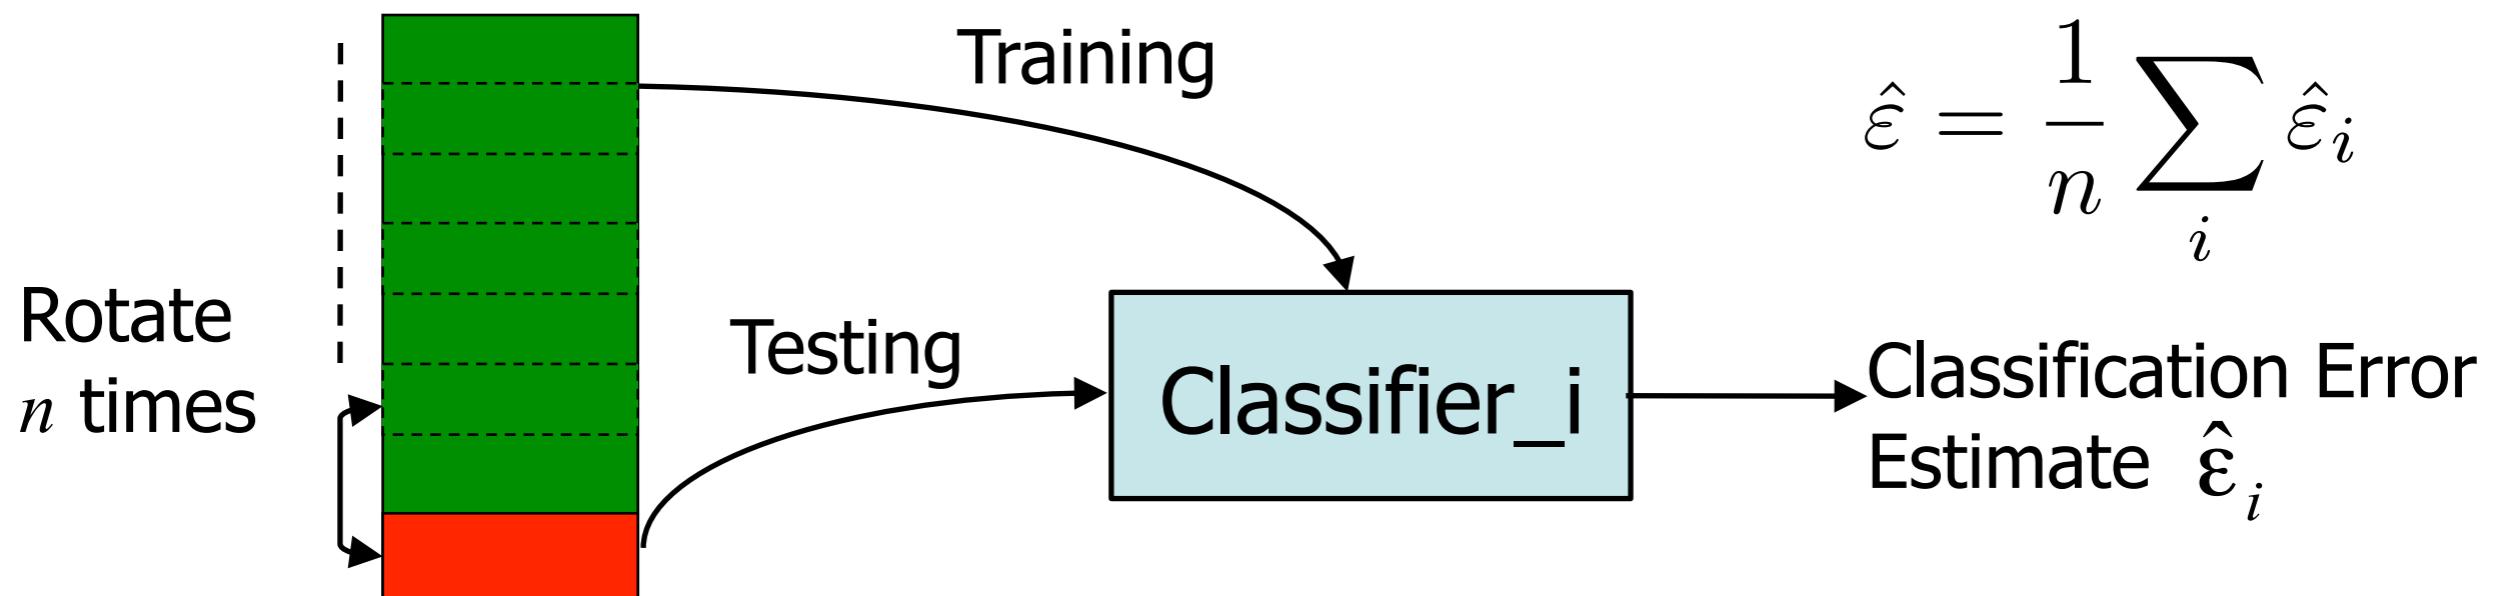
Zoom in of Fitbit data

- Ignore the time-stamp: (bug in the hardware??)



(still need to remove the low-frequency signal, but that is the next step)

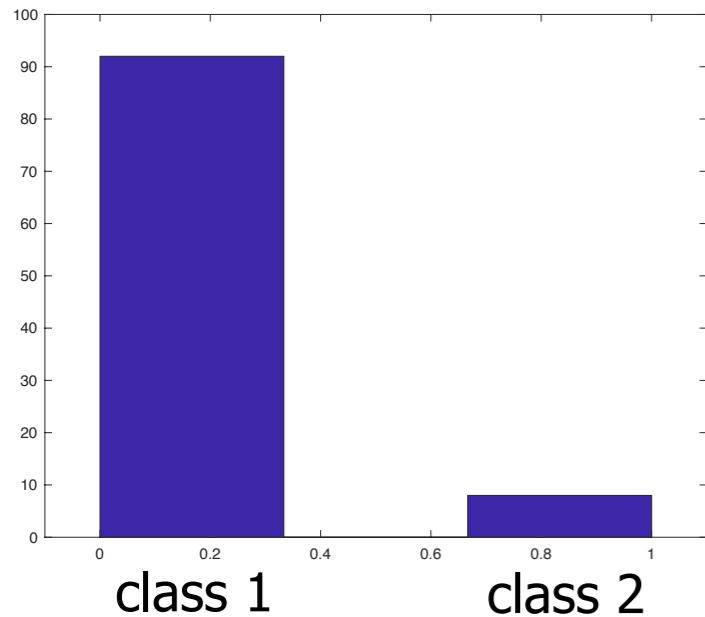
Cross-Validation



- Problems when classes are heavily imbalanced
- Problems when samples are very correlated

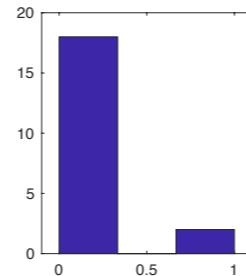
Standard cross-validation

full dataset:

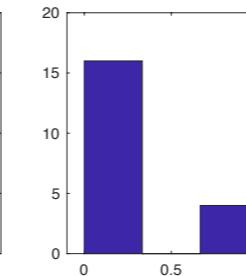


naive:

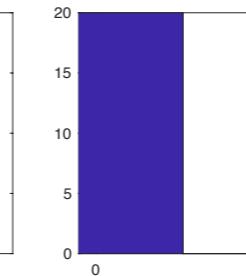
fold 1:



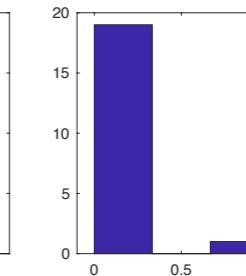
fold 2:



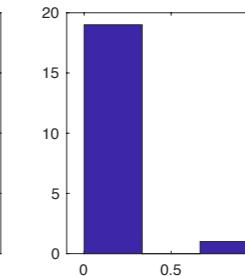
fold 3:



fold 4:



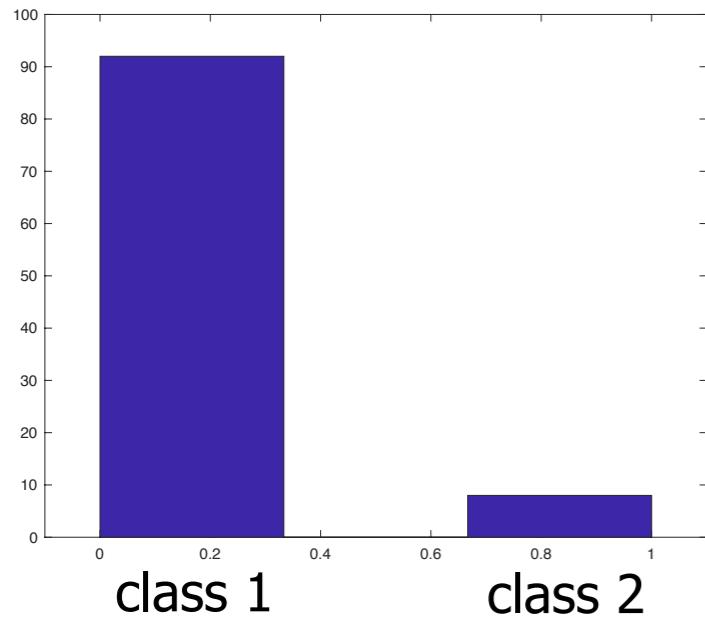
fold 5:



- Imbalanced classes: run the risk that the small class is lost in some of the folds

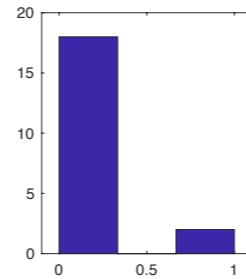
Stratified cross-validation

full dataset:

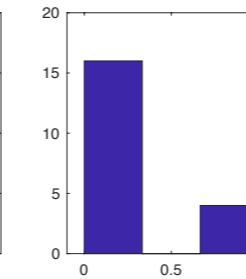


naive:

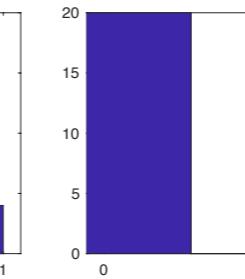
fold 1:



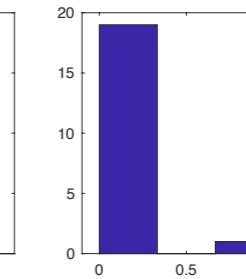
fold 2:



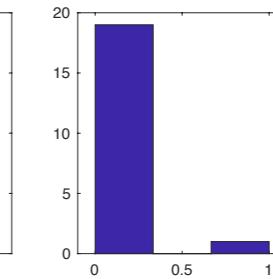
fold 3:



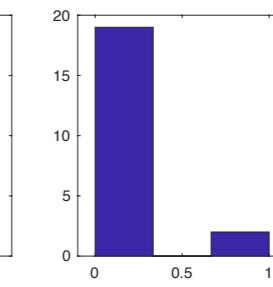
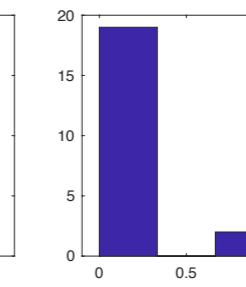
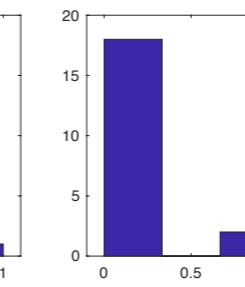
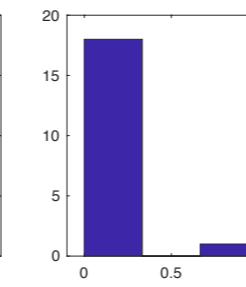
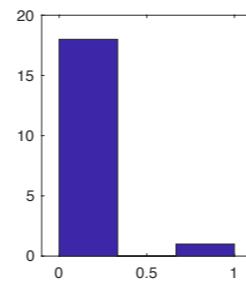
fold 4:



fold 5:



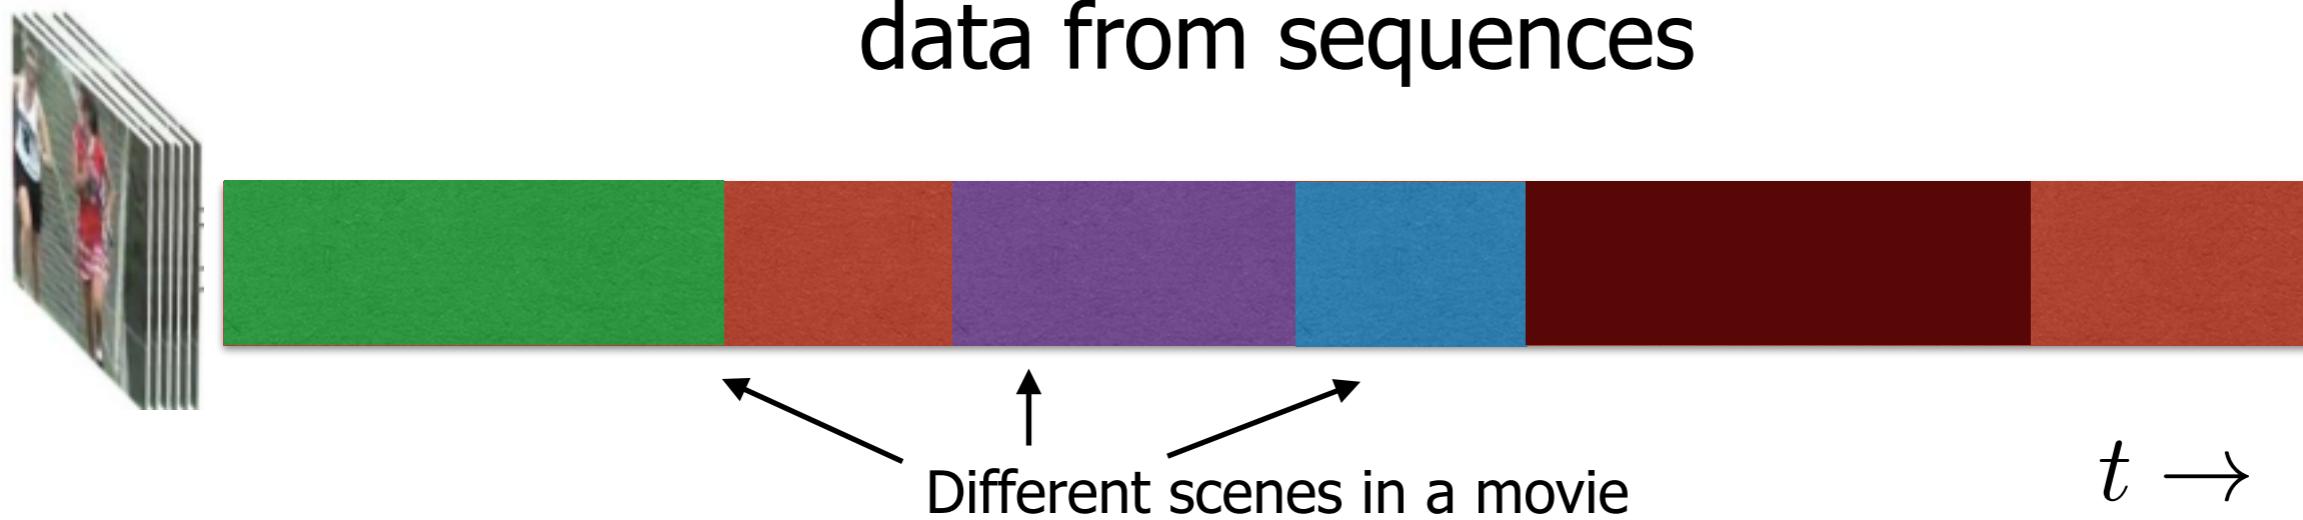
stratified:



- Imbalanced classes: run the risk that the small class is lost in some of the folds
- Try to get the same distribution in each fold:
 - 1. get the data from one class
 - 2. split in K folds
 - 3. combine the data from the different classes

Leave-one-group-out cross-validation

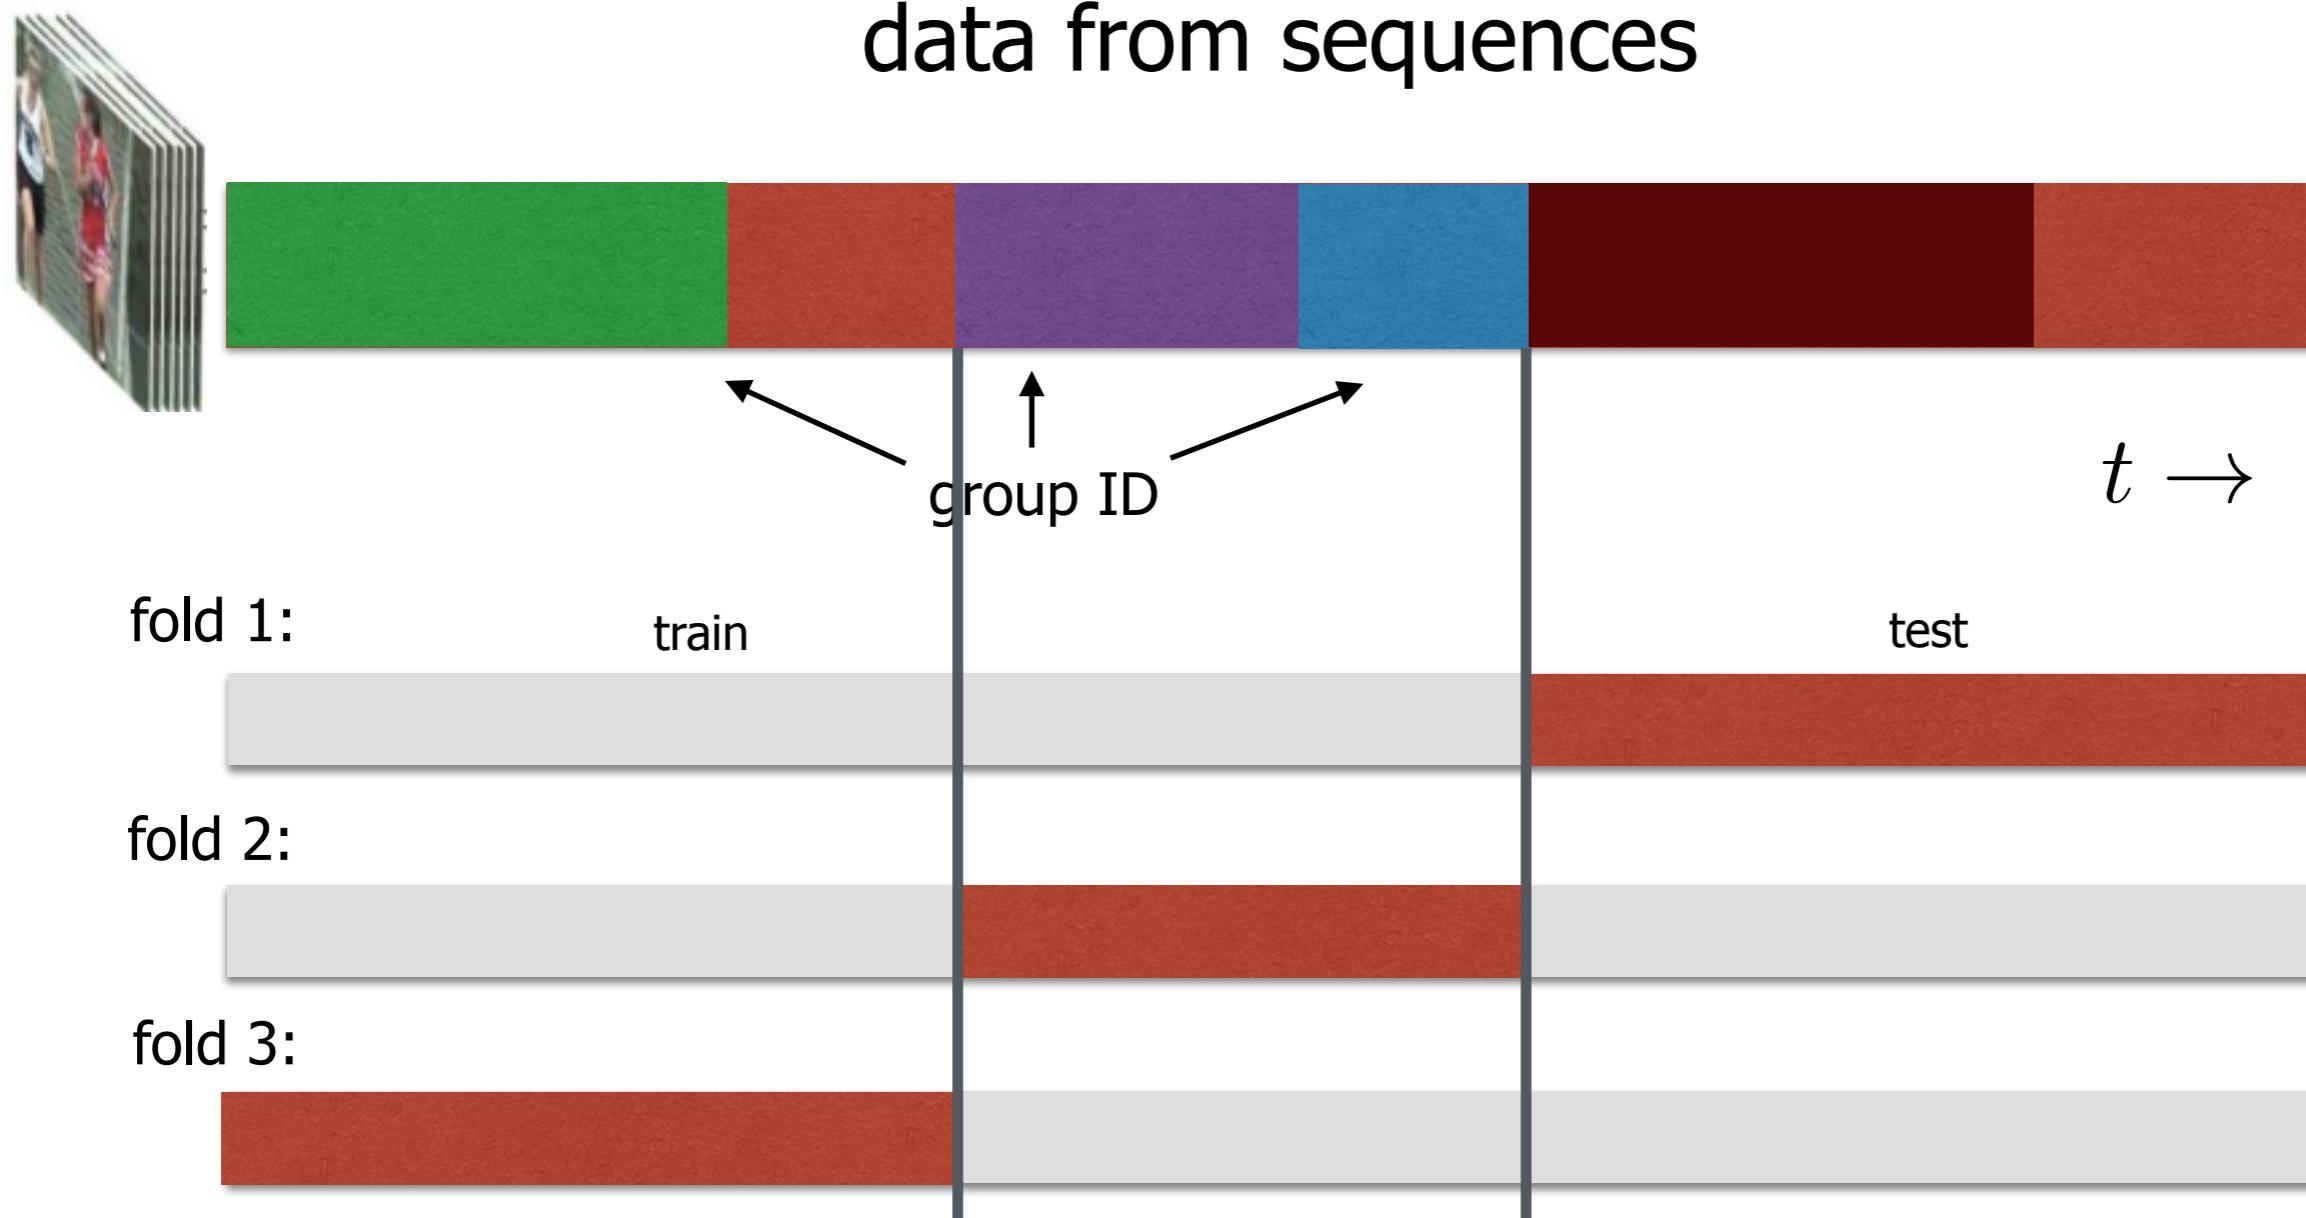
- For instance: data from different persons
data from different videos
data from sequences



- Samples within a group are heavily correlated
- Mixing them in training and testing gives (too) optimistic performance estimates

Leave-one-group-out cross-validation

- For instance: data from different persons
data from different videos
data from sequences



Leave-one-group-out cross-validation



- If you would use a random sample: will generalisation error increase/decrease/stay the same?

Comparing two methods

- Assume, two methods are evaluated using 10-fold cross-validation:

	fold 1	fold 2	fold 3	fold 4	fold 5	fold 6	fold 7	fold 8	fold 9	fold 10
Method 1	2.1	19.1	11.0	4.2	10.6	2.8	12.9	9.1	10.2	11.8
Method 2	4.0	19.3	10.6	5.1	10.8	4.2	12.1	11.0	11.5	11.7

- Which method is better?

Paired sample t-test

- Assume, two methods are evaluated using 10-fold cross-validation:

	fold 1	fold 2	fold 3	fold 4	fold 5	fold 6	fold 7	fold 8	fold 9	fold 10
Method 1	2.1	19.1	11.0	4.2	10.6	2.8	12.9	9.1	10.2	11.8
Method 2	4.0	19.3	10.6	5.1	10.8	4.2	12.1	11.0	11.5	11.7

- Which method is better?

mean (std)

Method 1 9.4 (5.2)

Method 2 10.0 (4.6)

- Is this significant?

Paired sample t-test

- Assume, two methods are evaluated using 10-fold cross-validation:

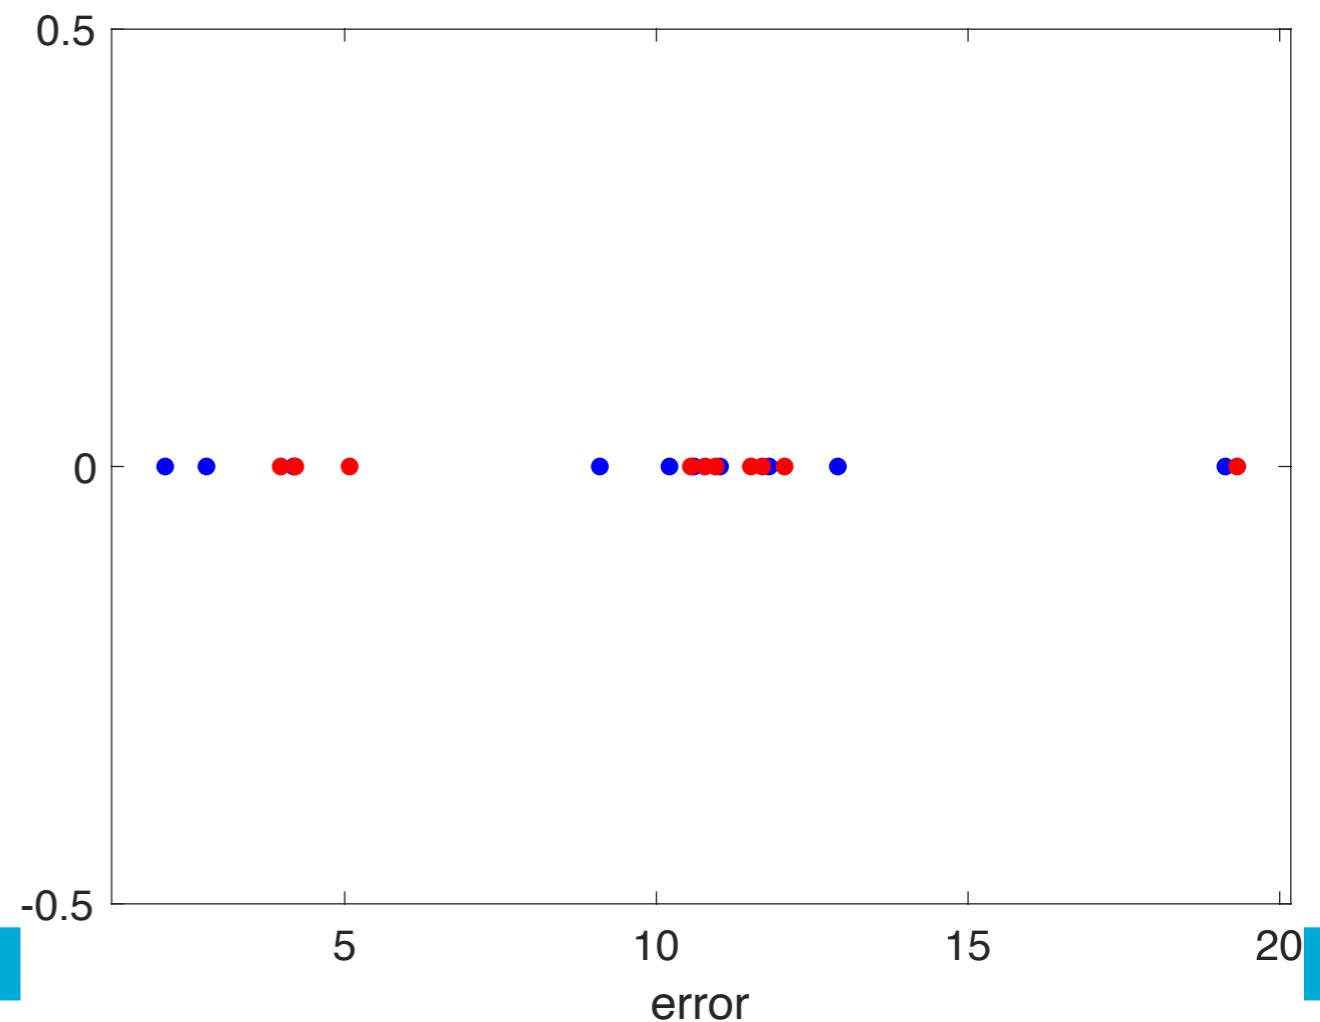
	fold 1	fold 2	fold 3	fold 4	fold 5	fold 6	fold 7	fold 8	fold 9	fold 10
Method 1	2.1	19.1	11.0	4.2	10.6	2.8	12.9	9.1	10.2	11.8
Method 2	4.0	19.3	10.6	5.1	10.8	4.2	12.1	11.0	11.5	11.7

- Which method is better?

mean (std)

Method 1	9.4 (5.2)
Method 2	10.0 (4.6)

- Is this significant?



Paired sample t-test

- If results come from the same data (identical folds):
look at the error **differences**:

	fold 1	fold 2	fold 3	fold 4	fold 5	fold 6	fold 7	fold 8	fold 9	fold 10
Difference	-1.9	-0.2	0.5	-0.9	-0.2	-1.4	0.9	-1.9	-1.3	0.1

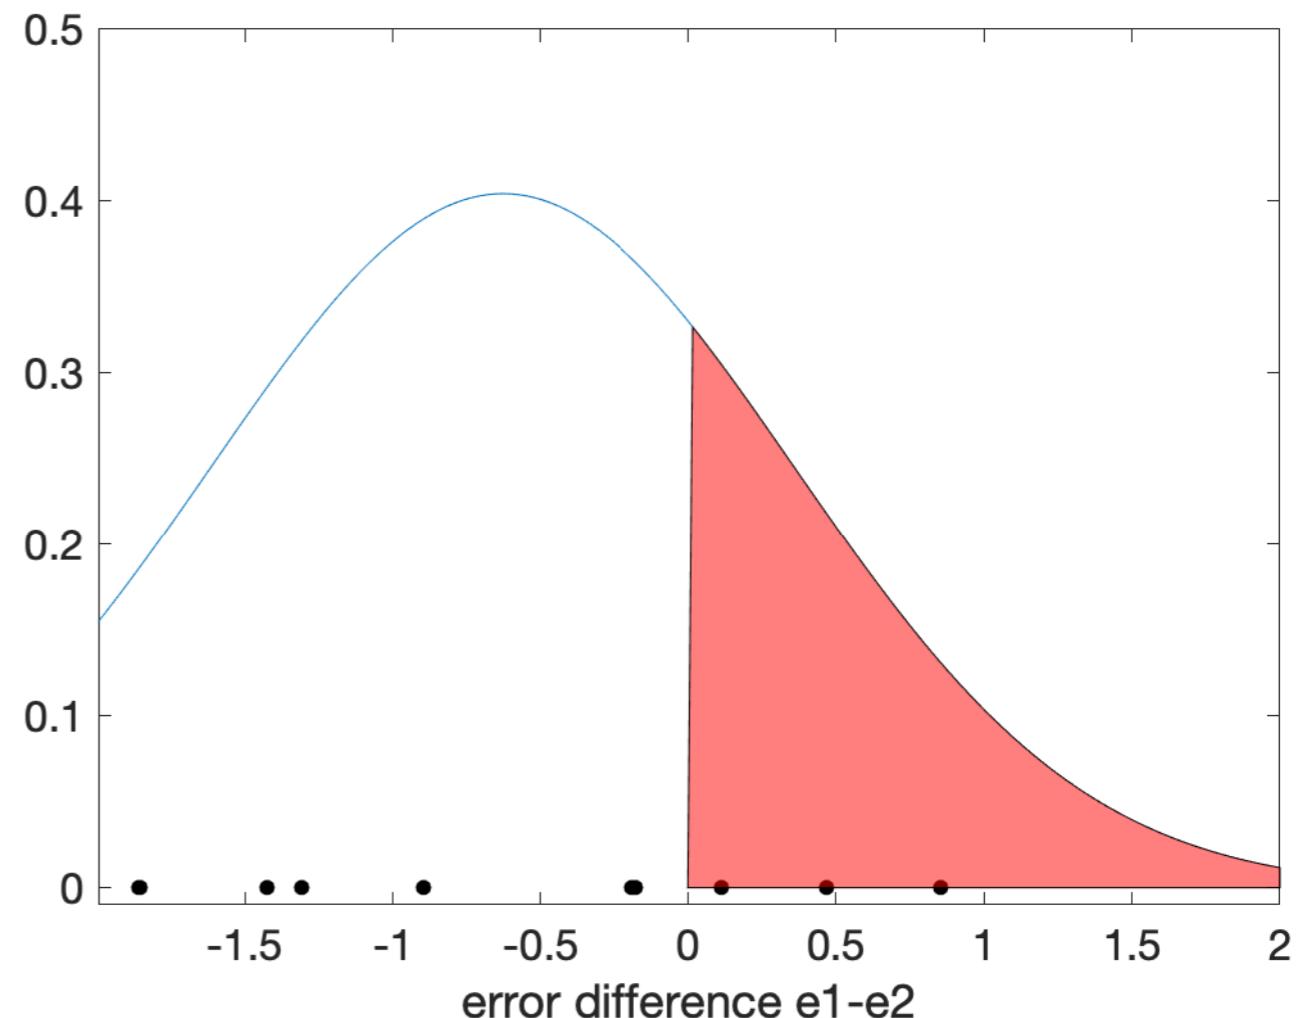
- Are these values significantly different from 0?

Paired sample t-test

- If results come from the same data (identical folds):
look at the error **differences**:

	fold 1	fold 2	fold 3	fold 4	fold 5	fold 6	fold 7	fold 8	fold 9	fold 10
Difference	-1.9	-0.2	0.5	-0.9	-0.2	-1.4	0.9	-1.9	-1.3	0.1

- Are these values significantly different from 0?



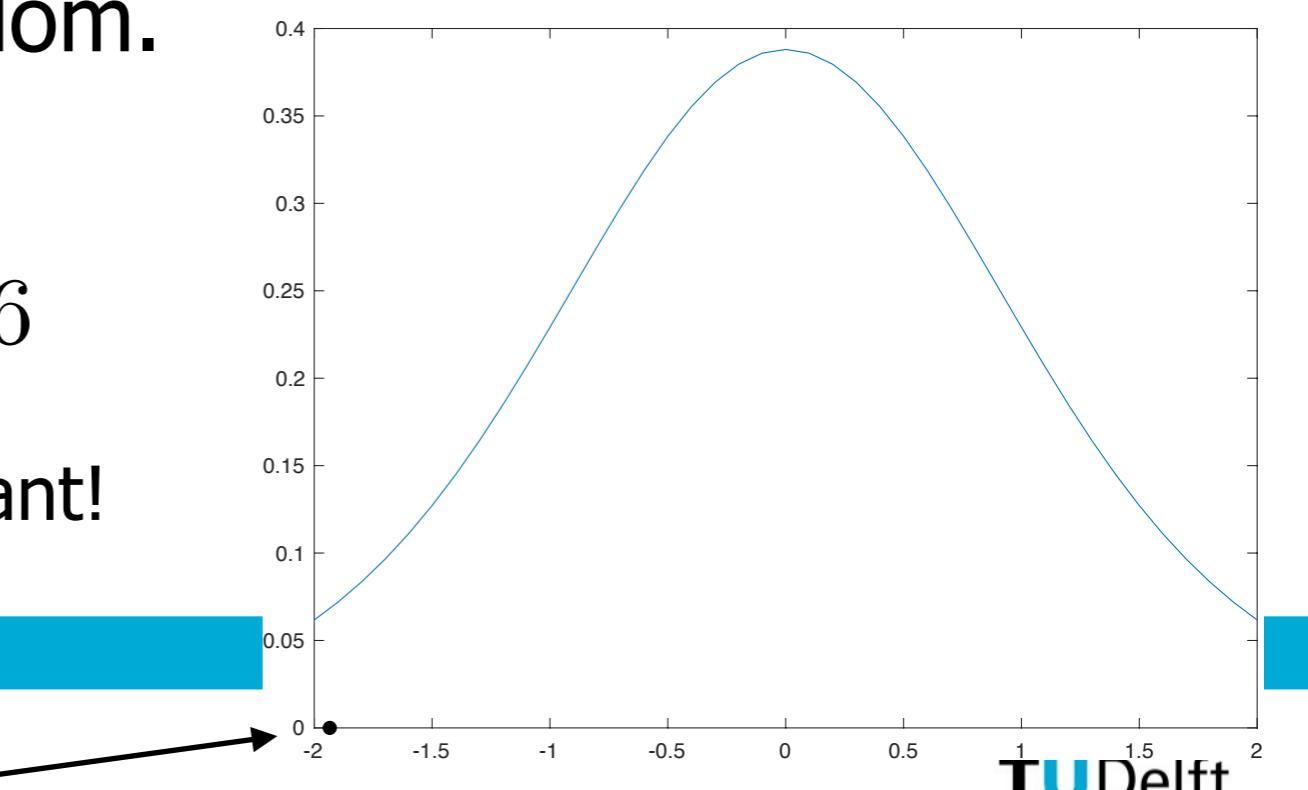
Paired sample t-test

- Are the averaged errors the same?
=
- Need a test-statistic
- You can show that the variable has a Student-t distribution with $(k-1)$ degrees of freedom.
- For us: $T=-1.93$

$$P(T \leq -1.93) = 0.0426$$

Just significant!

$$T = \frac{\bar{e}_1 - \bar{e}_2}{\sigma_{e_1 - e_2} / \sqrt{k}}$$

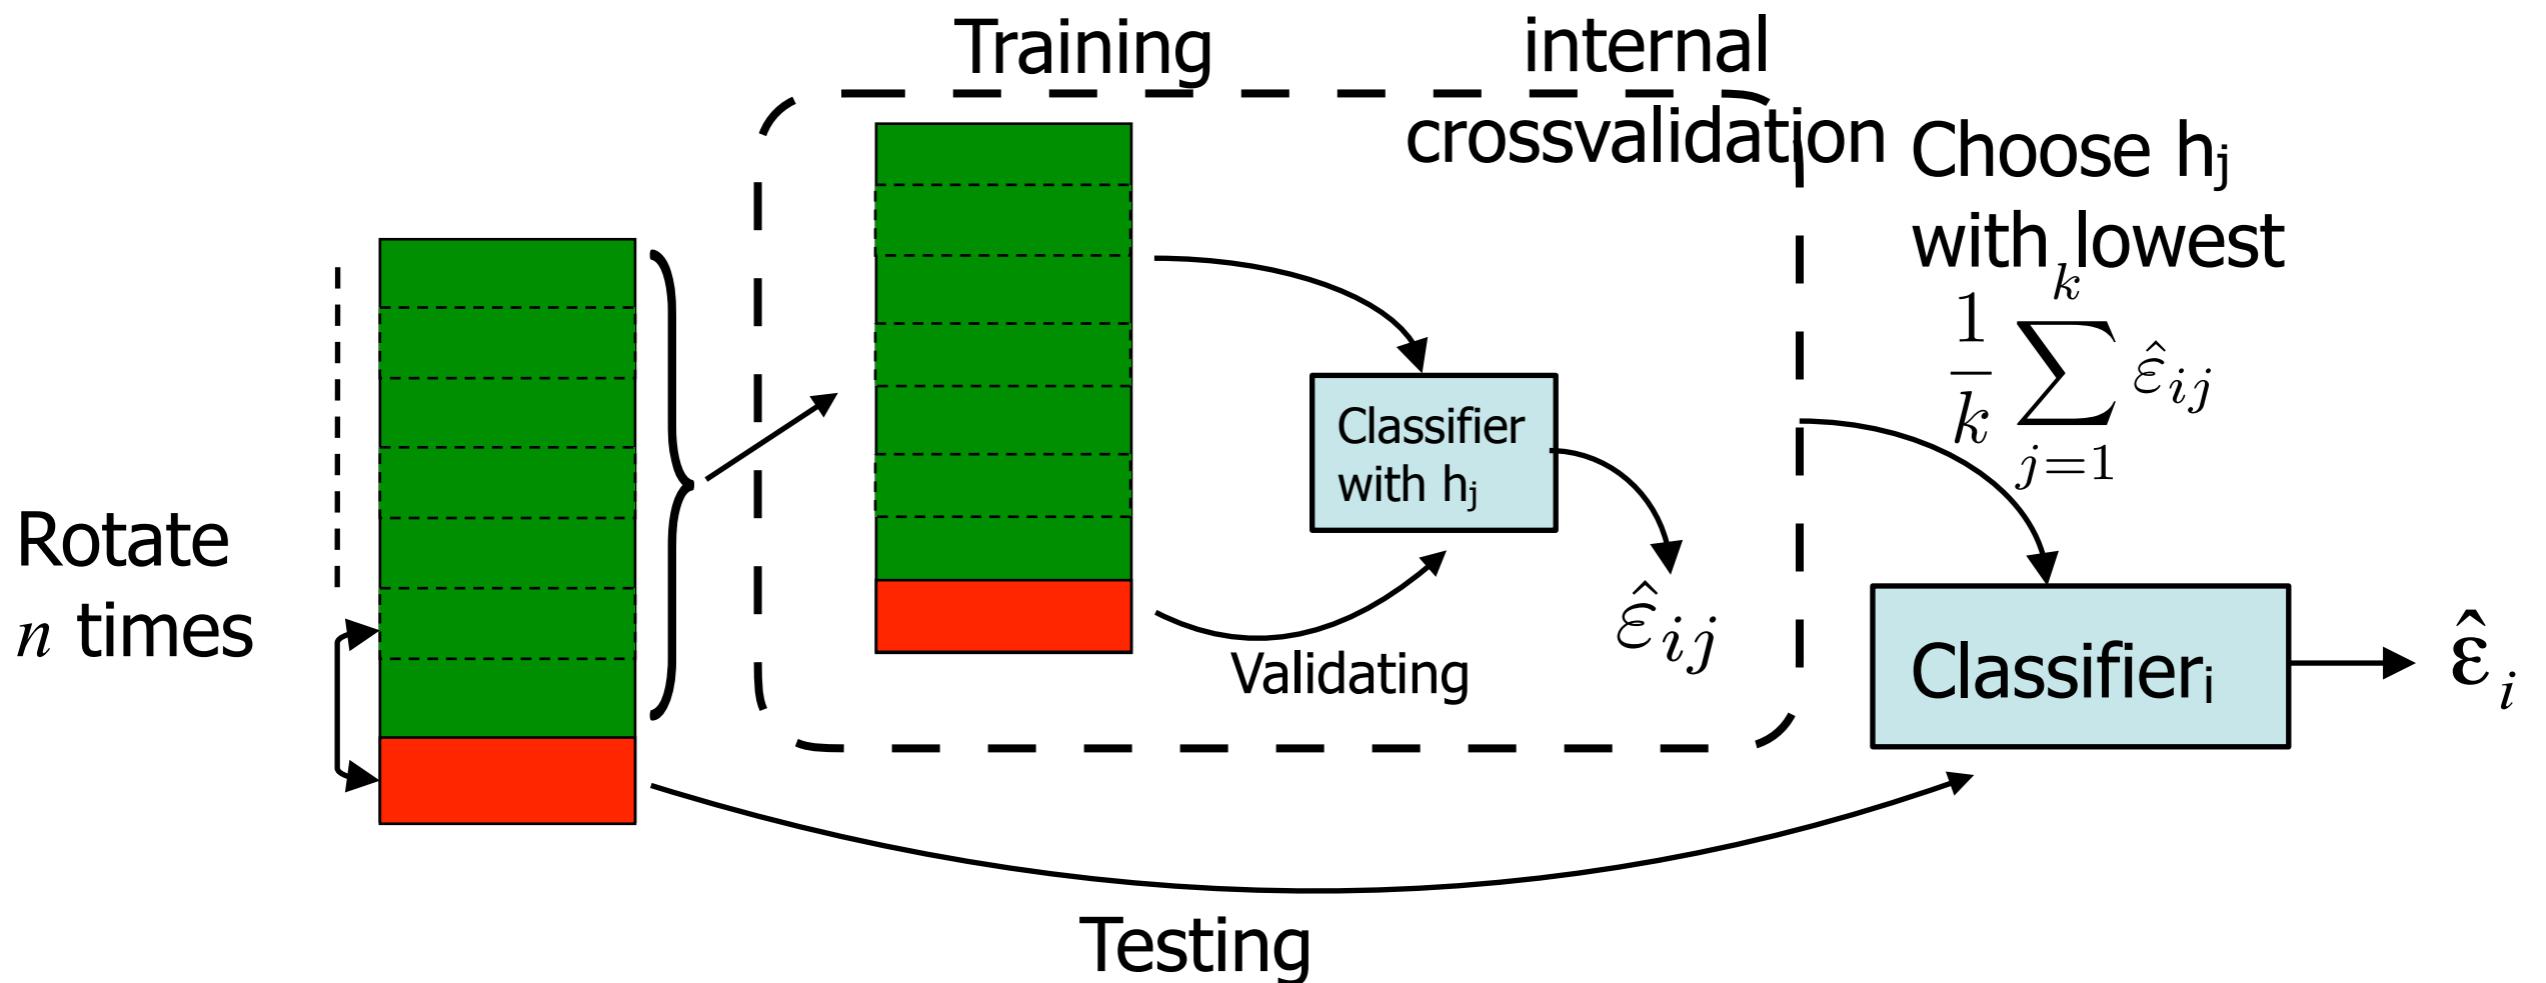


Optimisation of hyperparameters

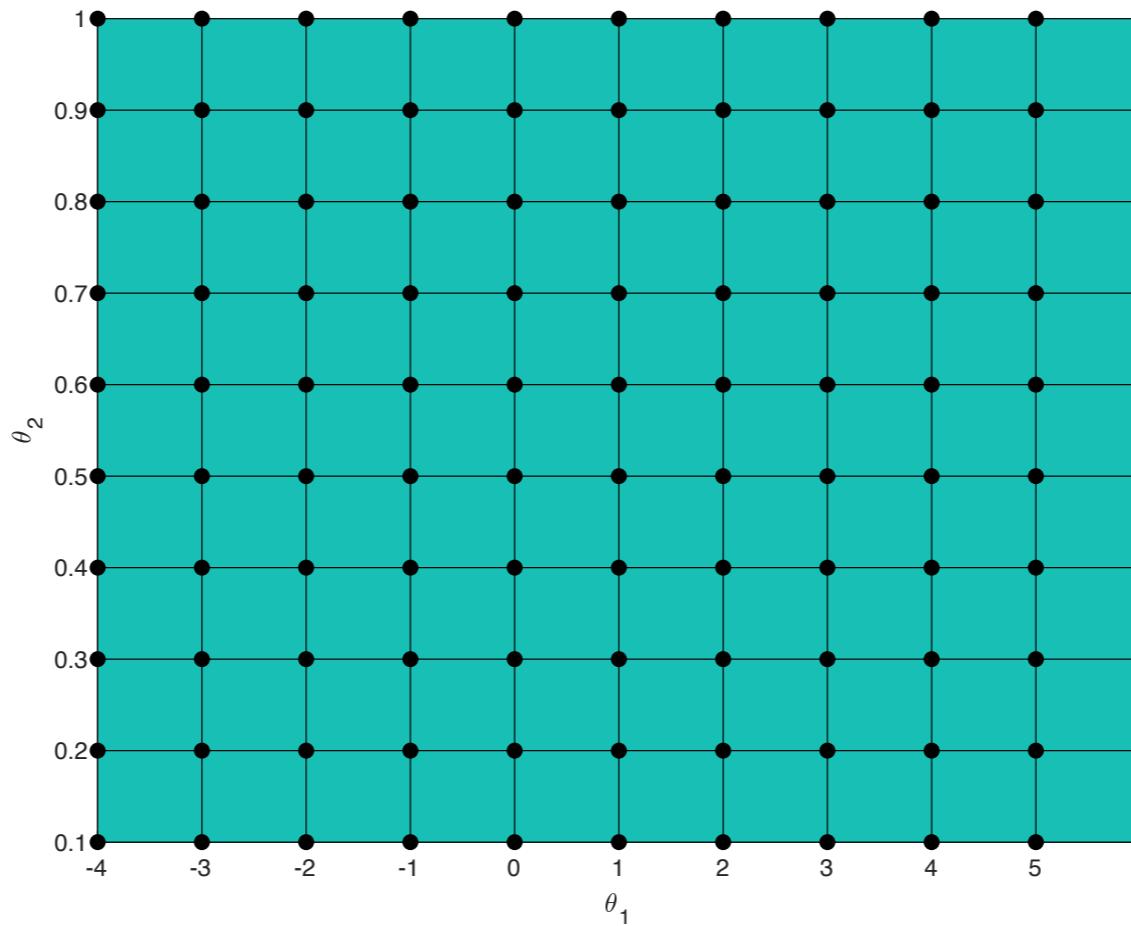
- Machine learning methods often have 'hyperparameters'
- Parzen density estimator: width parameter h
- k -nearest neighbour: number of neighbours k
- Decision trees: pruning method, stopping criterion
- Neural networks: architecture, learning rate, initialisation, batch size, regularisation, optimiser, ...
- DON'T optimise these numbers by looking at the test set!
Then you're **CHEATING!**

Double cross-validation

- To optimise over the hyperparameter $\{h_1, \dots, h_M\}$, do cross-validation **inside** another cross-validation:

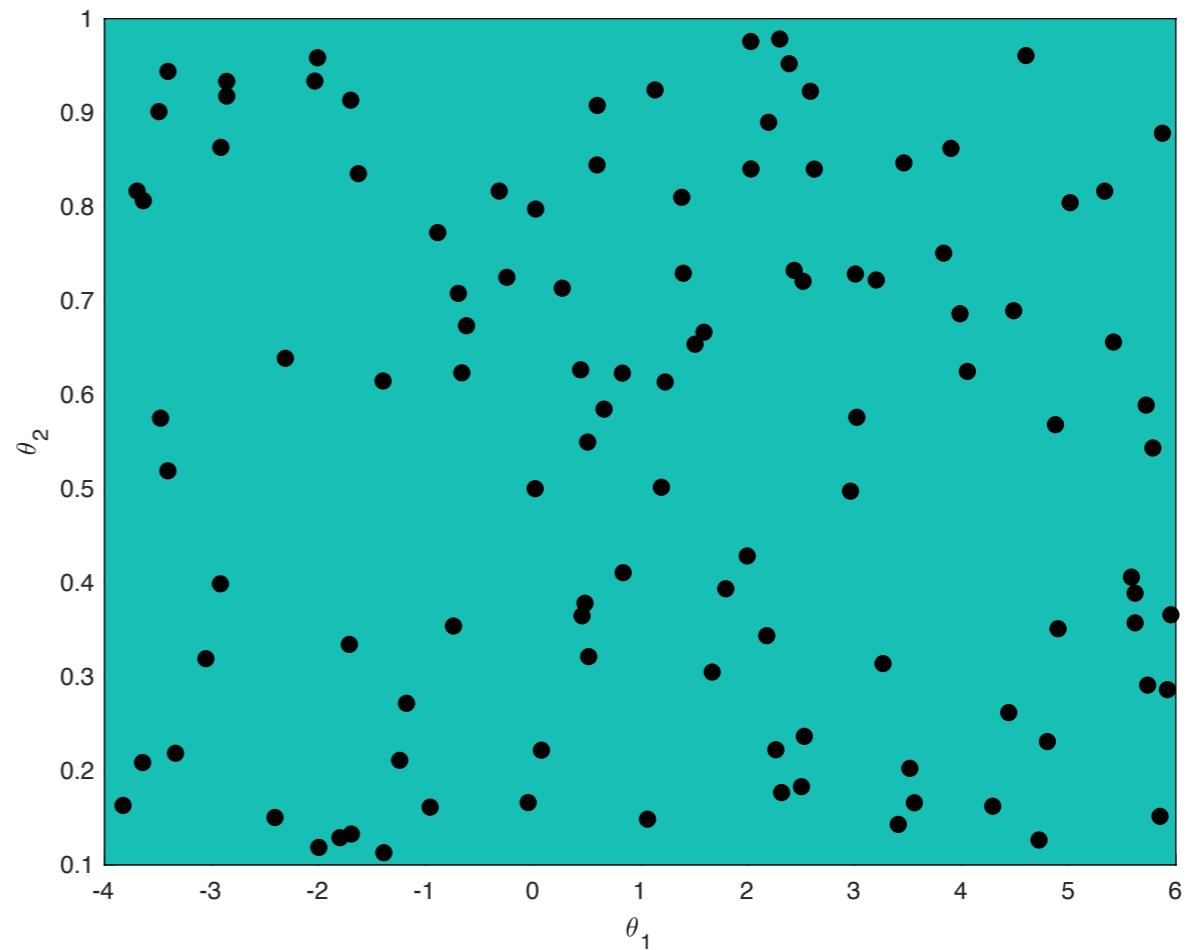
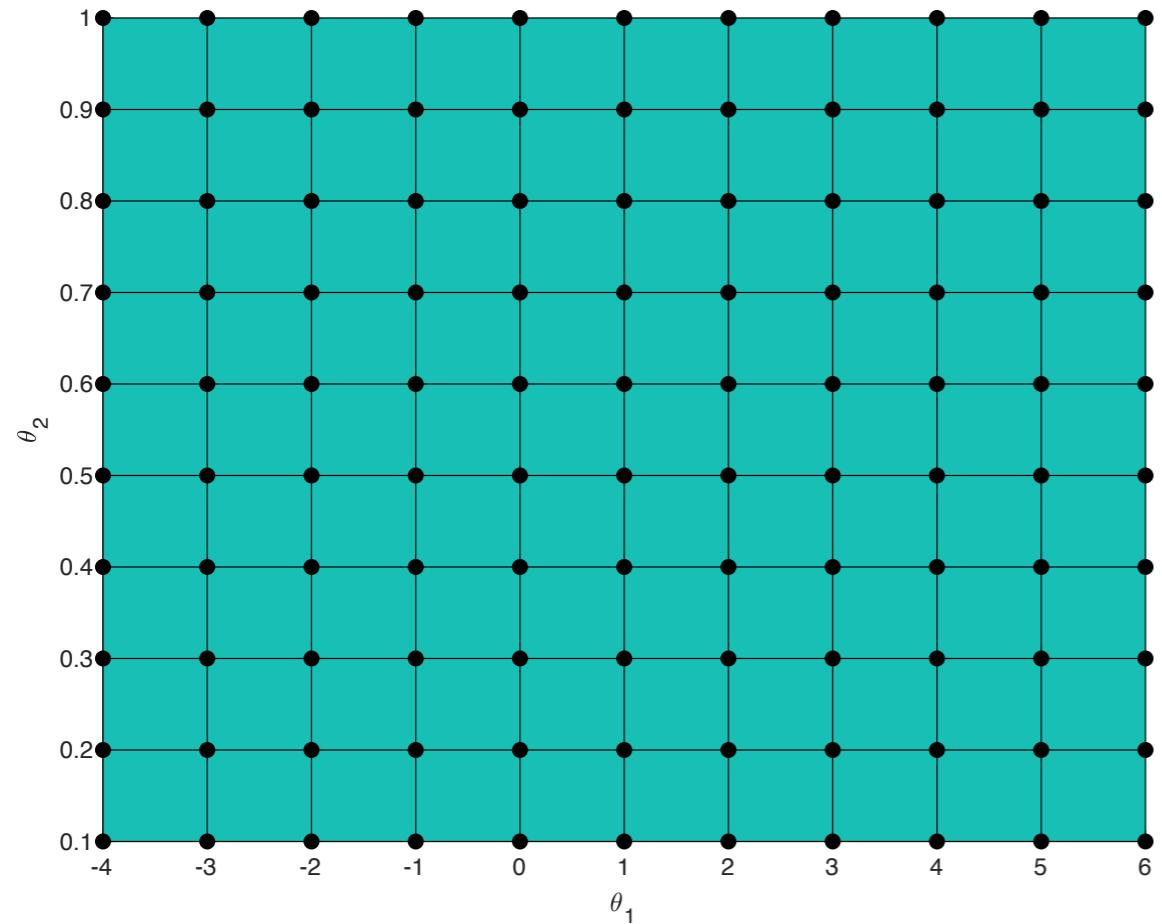


Optimisation of hyperparameters



- Grid search for 1 hyperparameter is fine
- For 2 parameters: Ok-ish, still do-able
- More than 2: Bayesian optimisation

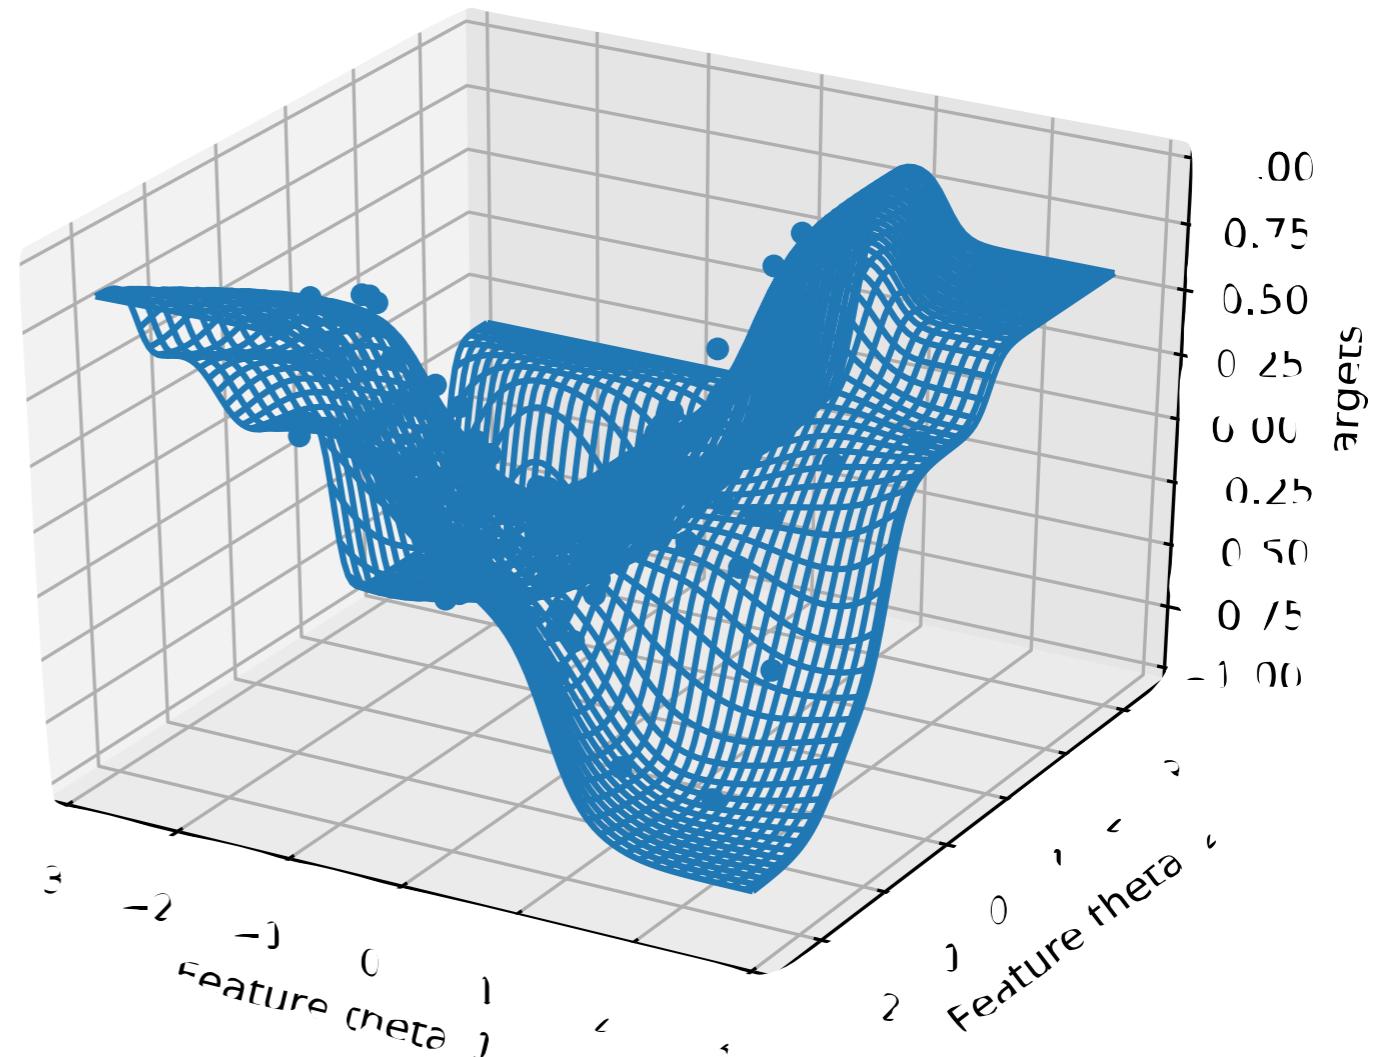
Optimisation of hyperparameters



- More than 2: Bayesian optimisation
- Introduce more variability in the values of the hyperparameter $\{h_1, \dots, h_M\}$

Bayesian optimisation

- When a new (random) hyperparameter settings are evaluated:
- fit a Gaussian Process regression
- THEN: Find minimum of the loss
- OR: Find maximum of the uncertainty, and evaluate that set of hyperparameters



After finding the minimum...

- The minimum error often not the most interesting (although often this is the 'proof' in articles)
- Try to understand the advantages/distadvantages:
 - What errors are made? (inspect objects, inspect labels)
 - What classes are problematic? (confusion matrix)
 - Does adding training data help? (learning curve)
 - How robust is the model?

Reporting results

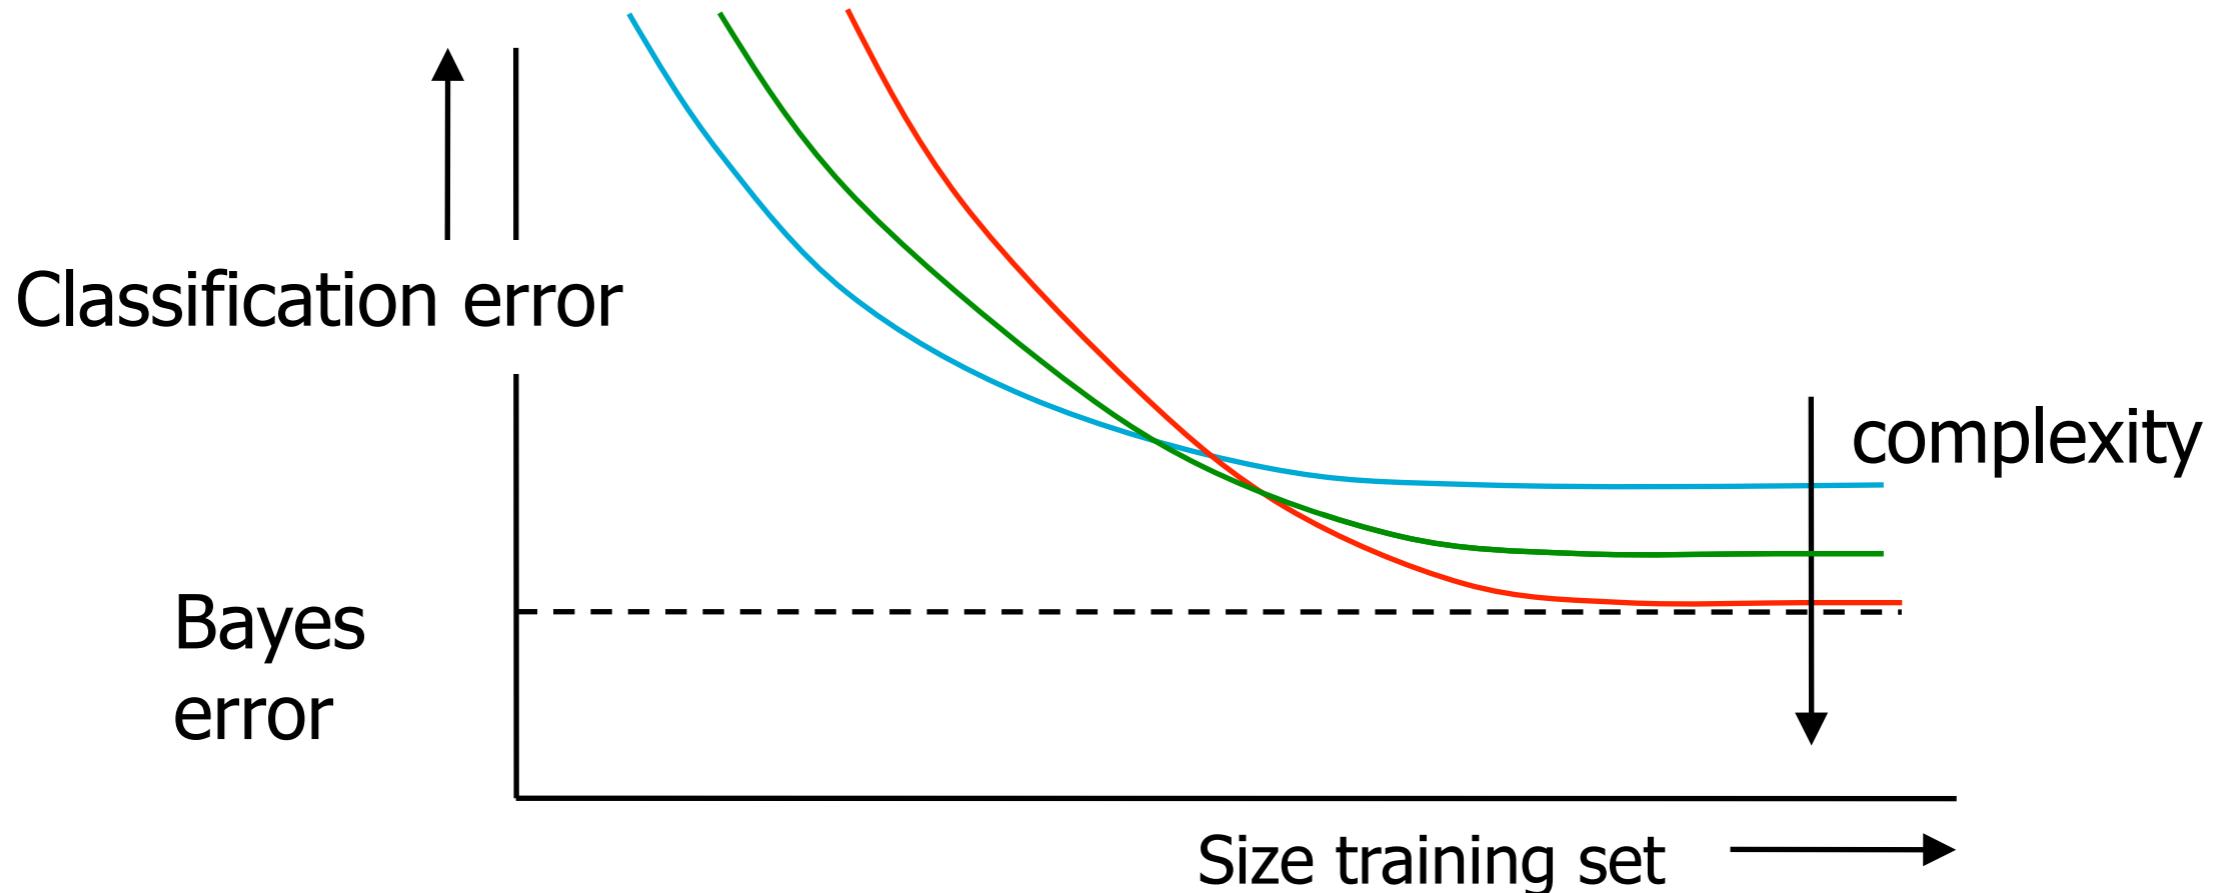
- Note, typically there's quite some noise involved in the experiments (split train-test, random initialisation,...)
- Experiments are typically repeated
- Do **NOT** just present a point estimate:

Method 1	9.4
Method 2	10.0
- But give some idea about standard deviations:

Method 1	9.4 (5.2)
Method 2	10.0 (4.6)

Different Classifier Complexity

- Don't claim to have the overall 'optimal' classifier
- Point out: what is good?
- What is bad?



Be critical!

- AI/ML/Deep Learning is a hype
- Do not trust all reported results
- Do not trust your own results

- 'Bold numbers' is not the goal of science:
know what the strengths and weaknesses of a model
are!
- If some method is very good, there should be
disadvantages

Conclusions

- Many possible Machine Learning methods, all with their strengths and weaknesses
- There is no overall best classifier
- For good generalisation, the bias of the model should fit the data (distribution)
- Still interested? Consider a Graduation project in ML
- Feedback on the course; wait for an upcoming Brightspace announcement