

Linear Regression

Jing Sun

Next Lecture – Nonlinear Regression by David :)

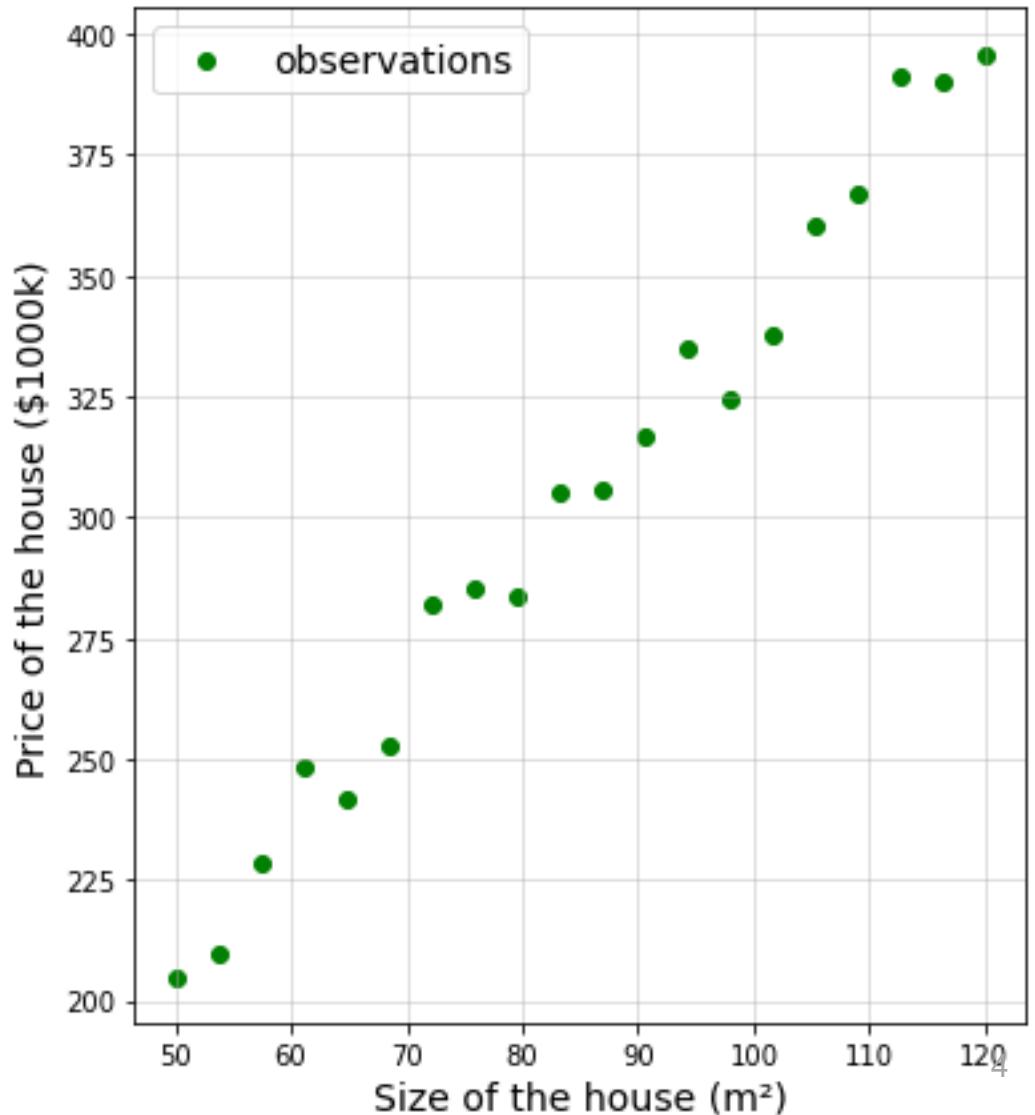
Agenda

- Linear Regression:
 - Ordinary Least Squares (OLS); R^2 ; Absolute Loss and Huber Loss.
- Overfitting and Bias-Variance Tradeoff.
- Regularizations:
 - Ridge a.k.a. L2 and Lasso a.k.a. L1.
- A little bit of the Classifiers:
 - (Multiclass) Logistic Regression and Perceptron (as there are some links).
- Bayesian Linear Regression (also prepare you for the next lecture).

OLS Linear Regression

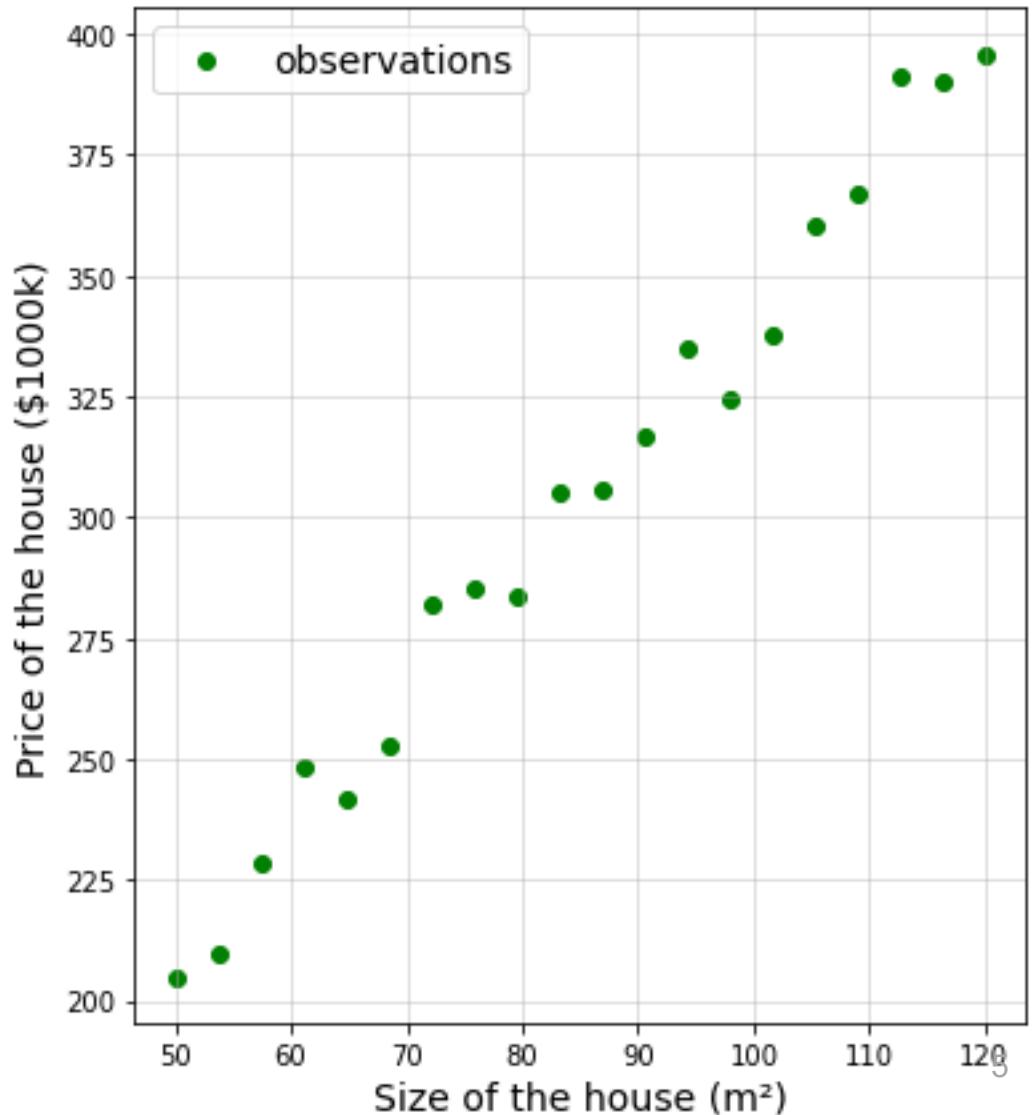
Linear Regression

- Given a set of features $x \in \mathbb{R}^d$,
we predict a response variable $y \in \mathbb{R}$
- Let's look at an example,
input (x-axis): size of the house
output (y-axis): price of the house



Linear Regression

- How to predict a new house's price given its size (an unseen x)?
- We can fit a linear function $f()$ to map x to y using the observations:



Linear Regression

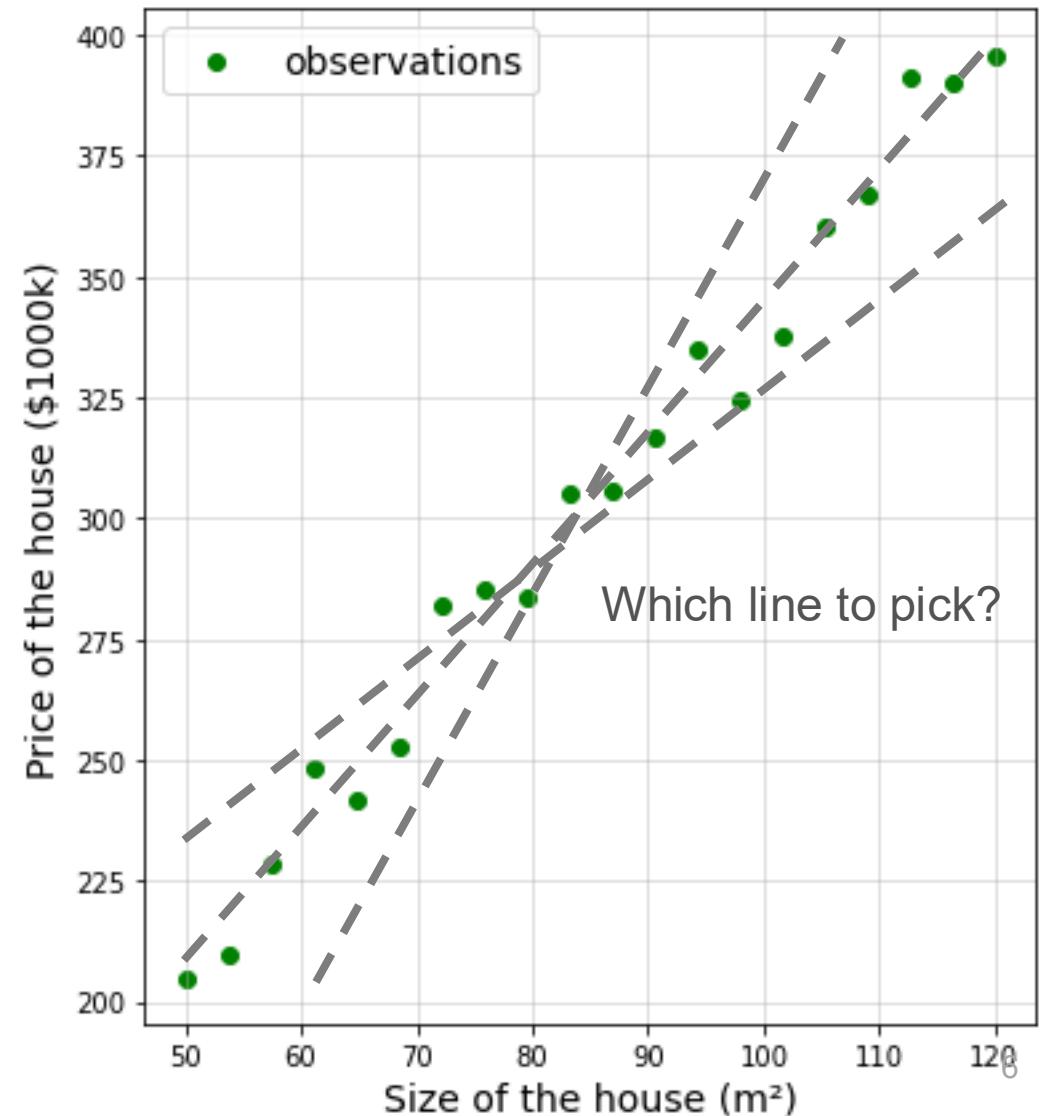
$$y = \beta_1 x_1 + \beta_0, \quad d = 1$$

slope intercept

$$\beta = \begin{bmatrix} \beta_0 \\ \beta_1 \\ \beta_2 \\ \vdots \\ \beta_d \end{bmatrix}, \quad x = \begin{bmatrix} 1 \\ x_1 \\ x_2 \\ \vdots \\ x_d \end{bmatrix}$$

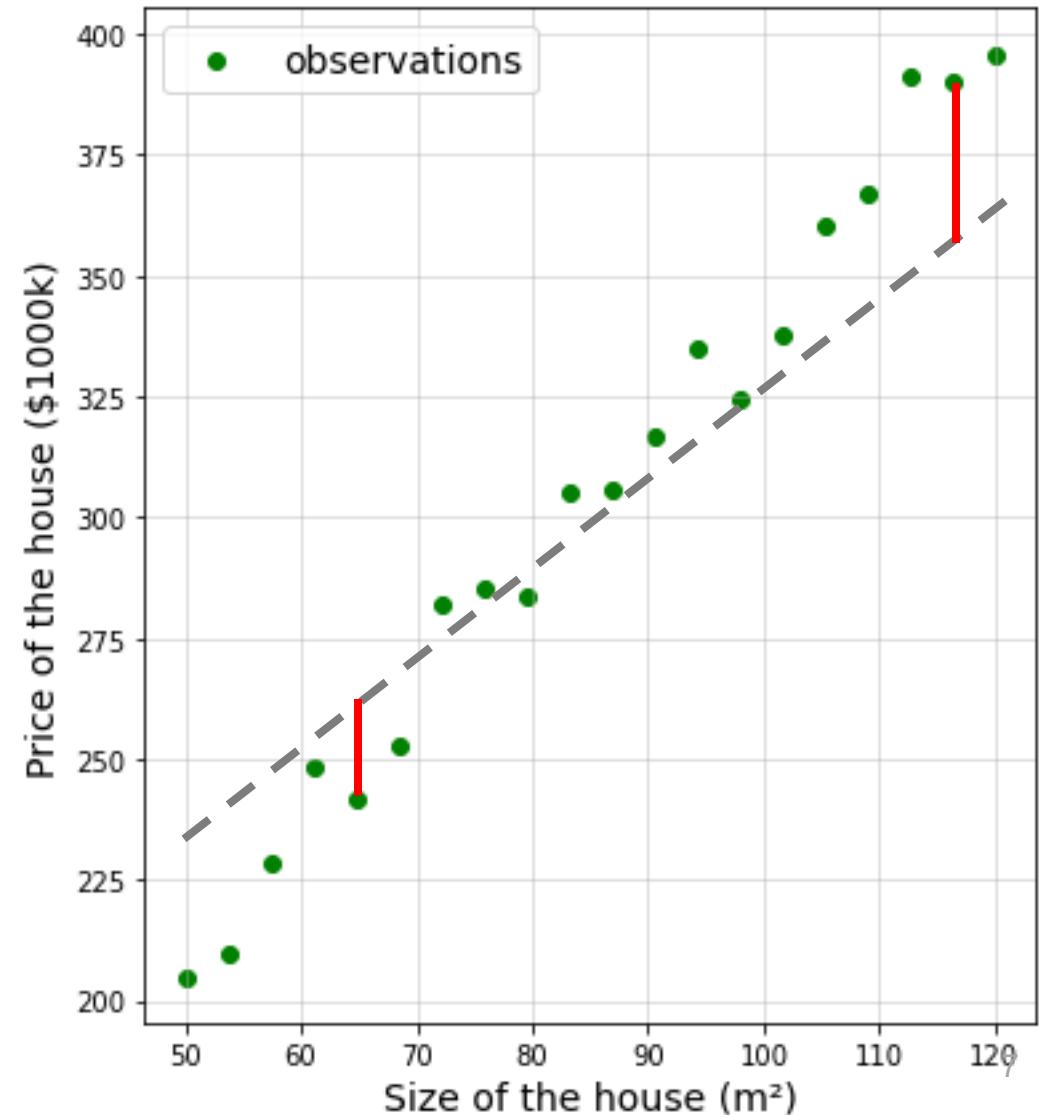
$$y = x^T \beta$$

- How to estimate β ?



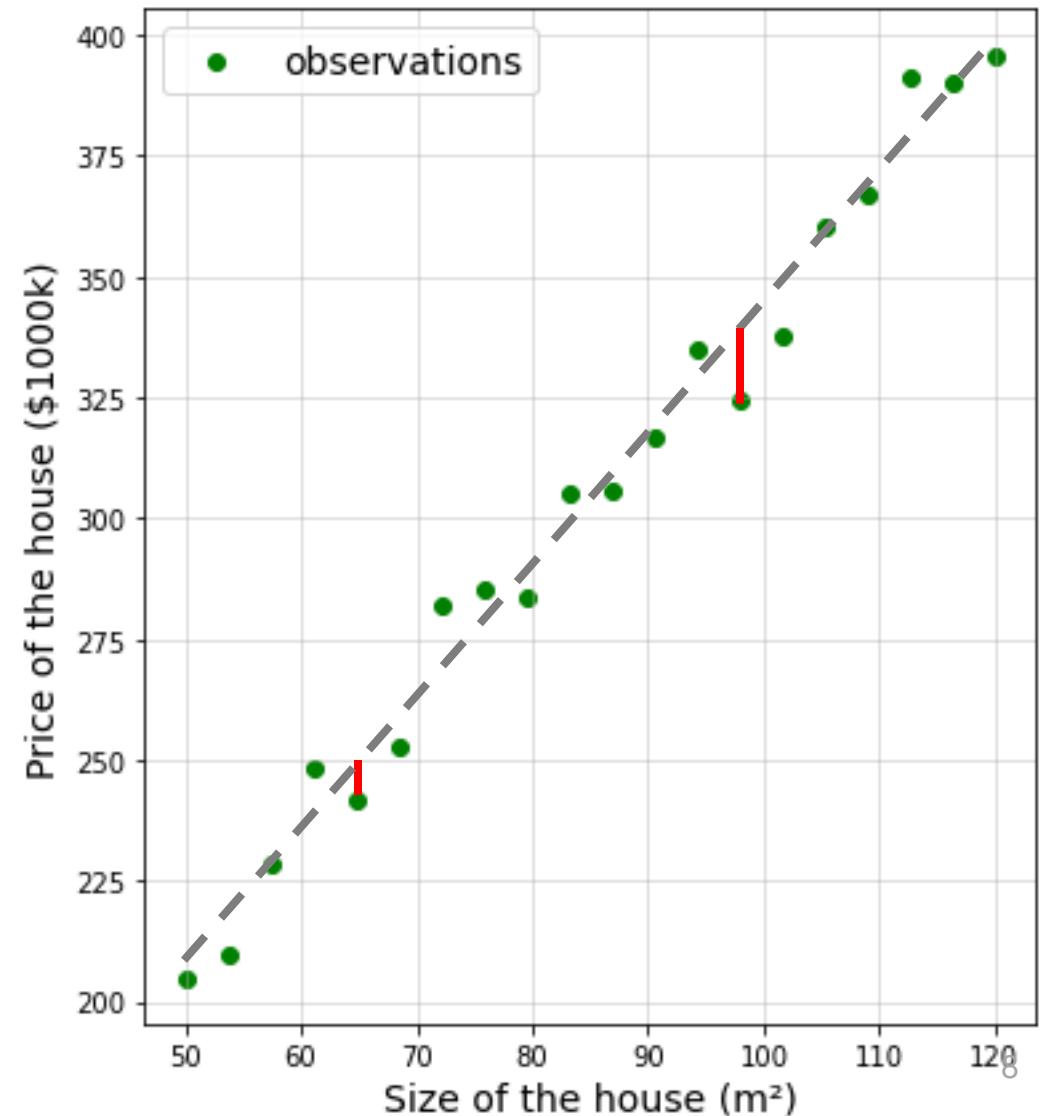
Residuals

- Well, we want β to make y and \hat{y} as close as possible, that is to say, to minimize the “residual”.
- What is residual r_n ?
- It is the **difference** between the **observed value** and the **predicted value** for a given data point.
- $r_n = y_n - \hat{y}_n = y_n - x_n^T \hat{\beta}$



OLS

- Minimize the residual sum of squares!
- $RSS = \sum_{n=1}^N (r_n)^2 = \sum_{n=1}^N (y_n - \hat{y}_n)^2$
- sensitive to outliers



OLS

for the intercept

- $X = \begin{bmatrix} 1 & x_{11} & x_{12} & x_{13} & \dots & x_{1d} \\ 1 & x_{21} & x_{22} & x_{23} & \dots & x_{2d} \\ 1 & \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & x_{N1} & x_{N2} & x_{N3} & \dots & x_{Nd} \end{bmatrix}, \quad \beta = \begin{bmatrix} \beta_0 \\ \beta_1 \\ \beta_2 \\ \vdots \\ \beta_d \end{bmatrix}, \quad y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_N \end{bmatrix},$

- Parameter estimation:

$$\hat{\beta}_{OLS} = \underset{\beta \in \mathbb{R}^{d+1}}{\operatorname{argmin}} \sum_{n=1}^N (y_n - x_n^T \beta)^2 = \underset{\beta \in \mathbb{R}^{d+1}}{\operatorname{argmin}} (y - X\beta)^T (y - X\beta)$$

- Solution given at

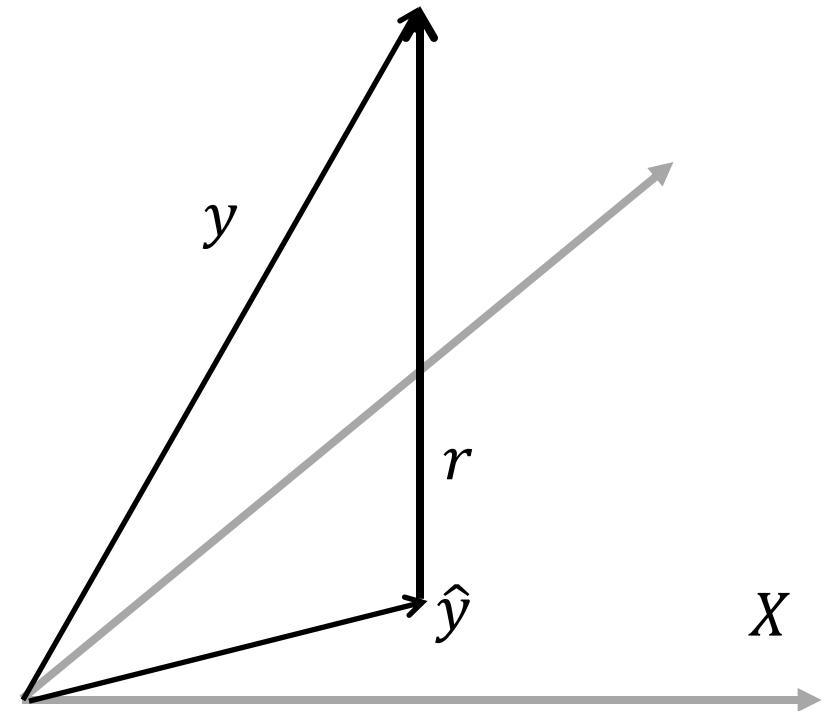
$$\frac{\partial}{\partial \beta} (y - X\beta)^T (y - X\beta) = 0 \quad \rightarrow \quad \hat{\beta}_{OLS} = (X^T X)^{-1} X^T y$$

OLS – another angle

- Geometric Interpretation to OLS:
- Find the **orthogonal projection** of the y onto the d -dimensional subspace spanned by the X .

$$X^T(y - X\hat{\beta}_{OLS}) = 0$$

$$\rightarrow \hat{\beta}_{OLS} = (X^T X)^{-1} X^T y$$



R^2 : Coefficient of Determination

- The proportion of variability in Y that can be explained by the regression.

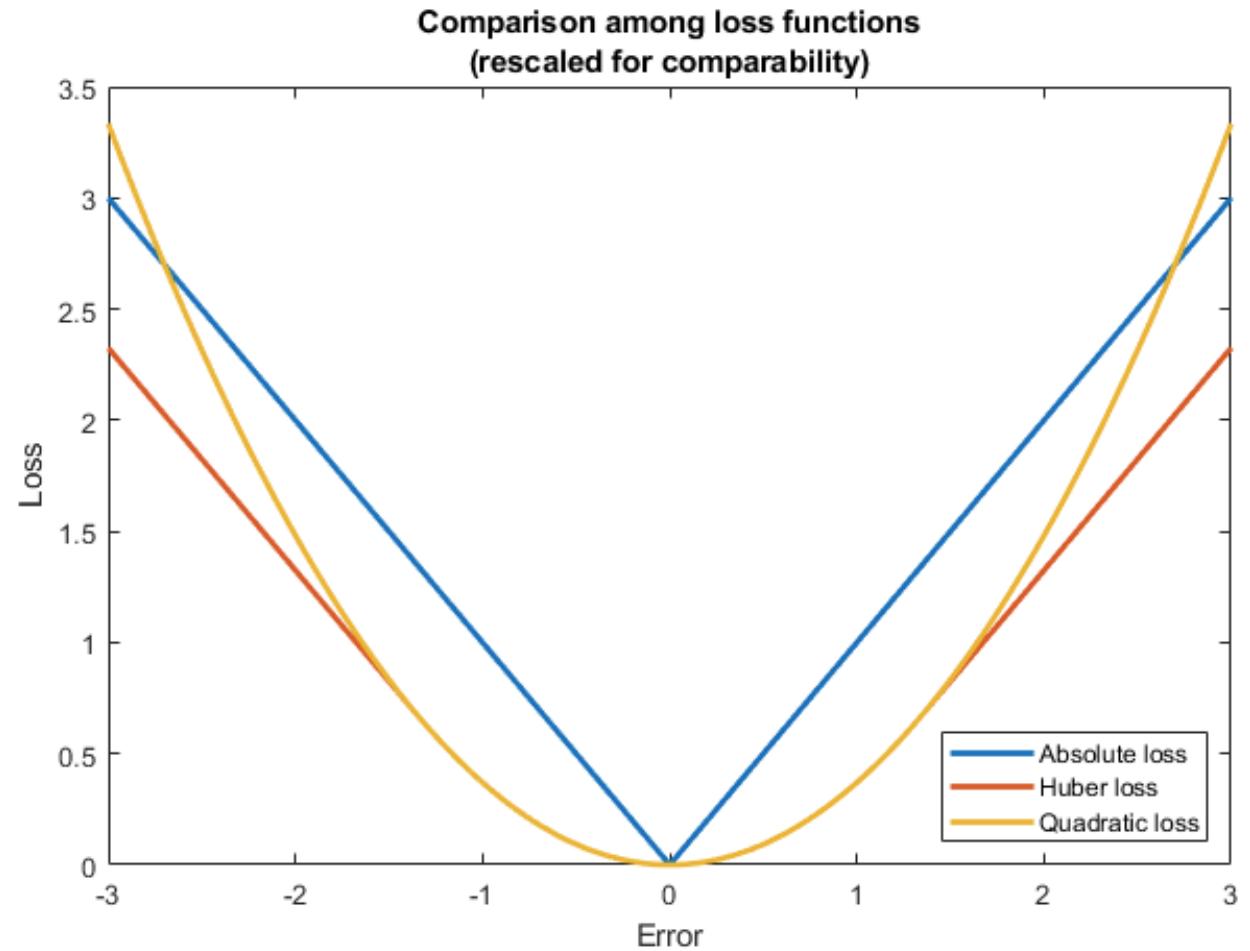
Total Sum of Squares $TSS = \sum (y_n - \bar{y})^2$

$$R^2 = \frac{TSS - RSS}{TSS} = 1 - \frac{RSS}{TSS}$$

$$R^2 = \frac{\text{variation}(data) - \text{variation}(fit)}{\text{variation}(data)}$$

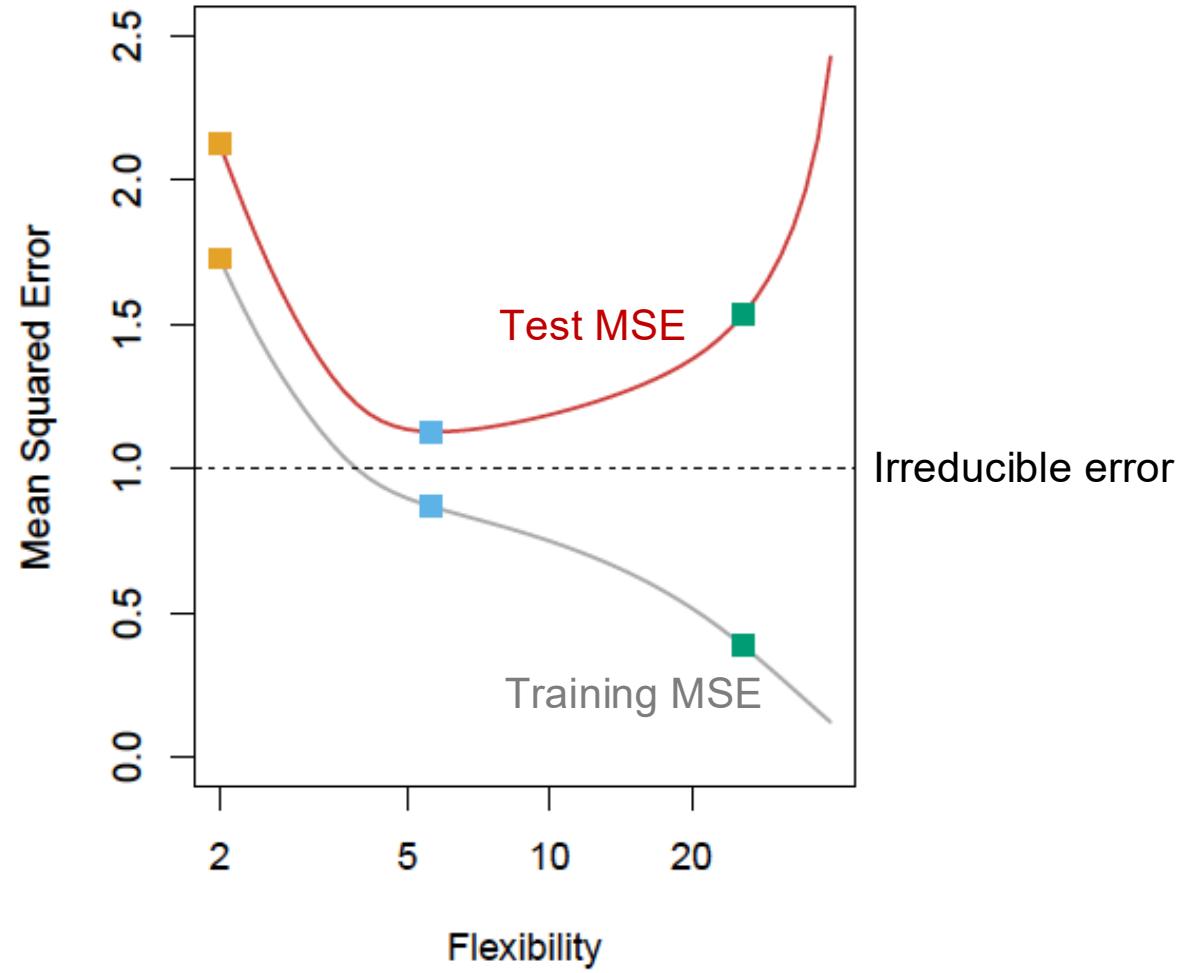
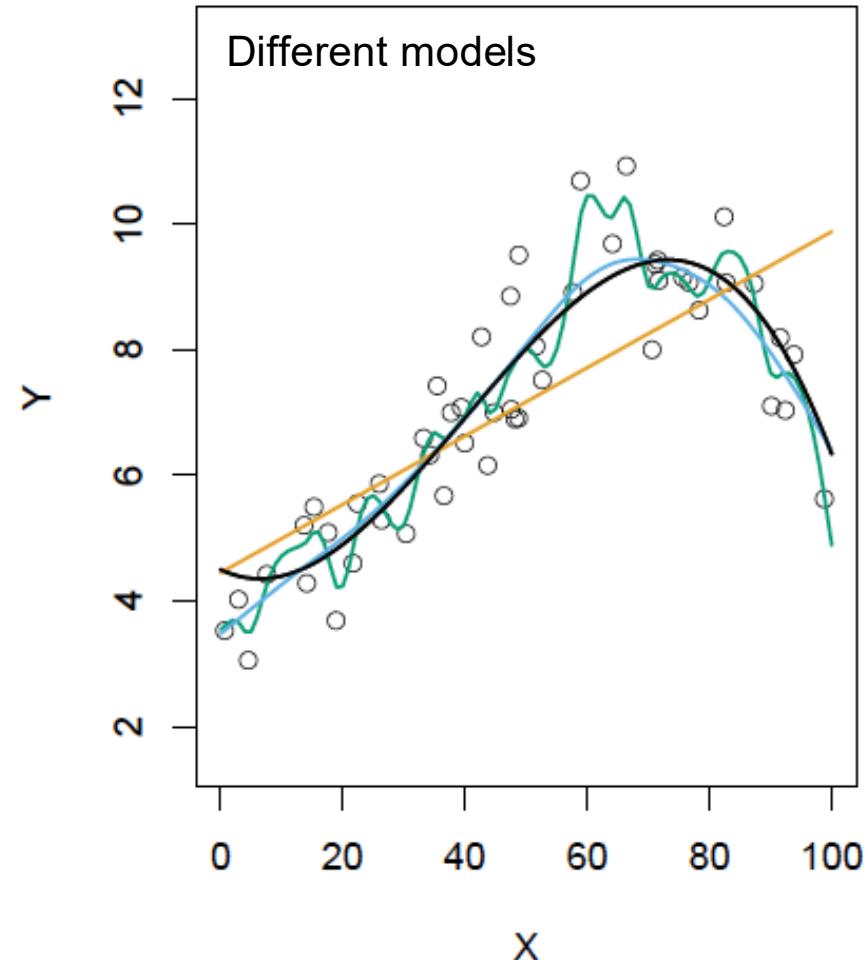
Other Loss Functions

- Absolute Loss= $|y - \hat{y}|$
- Huber Loss=
$$\begin{cases} \frac{1}{2}(y - \hat{y})^2, & \text{if } |y - \hat{y}| \leq \delta \\ \delta \left(|y - \hat{y}| - \frac{\delta}{2} \right), & \text{otherwise} \end{cases}$$
- Quadratic Loss: $(y - \hat{y})^2$



Overfitting and Bias-Variance Tradeoff

Bia-Variance Tradeoff



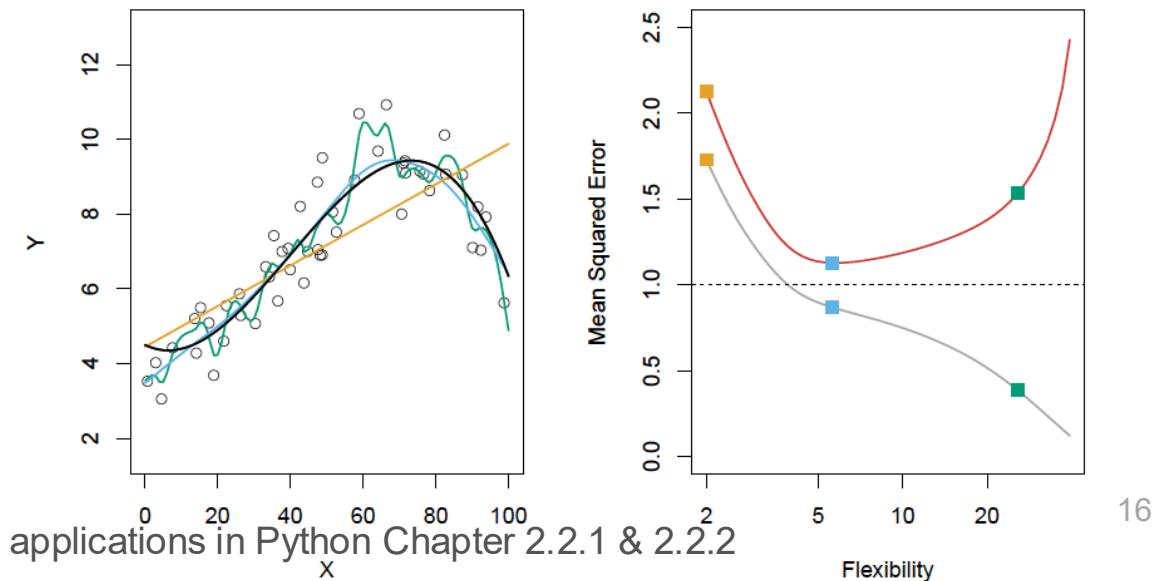
Bia-Variance Tradeoff

$$\begin{aligned} E(Y - \hat{Y}) &= E[f(X) + \epsilon - \hat{f}(X)]^2 \\ &= \underbrace{[f(X) - \hat{f}(X)]^2}_{= \text{reducible error}} + \text{Var}(\epsilon) \\ &= \text{reducible error} + \text{irreducible error} \\ &= \text{Bias}(\hat{f}(X))^2 + \text{Var}(\hat{f}(X)) + \text{Var}(\epsilon) \end{aligned}$$

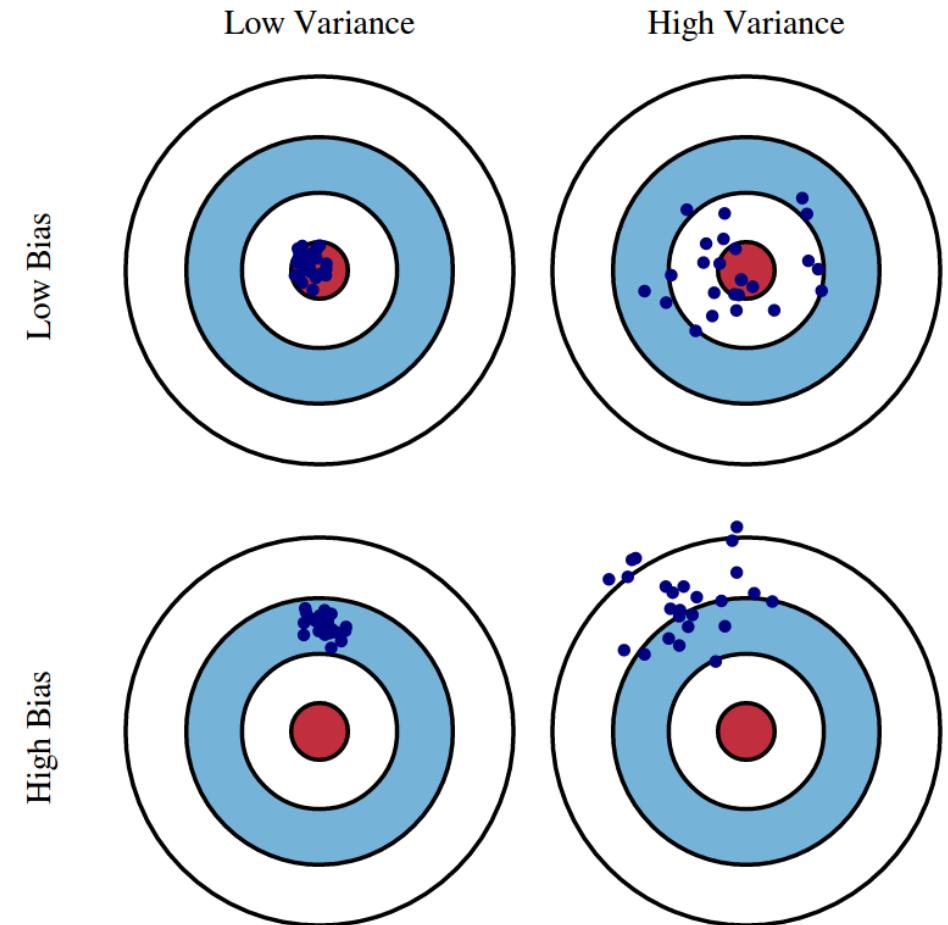
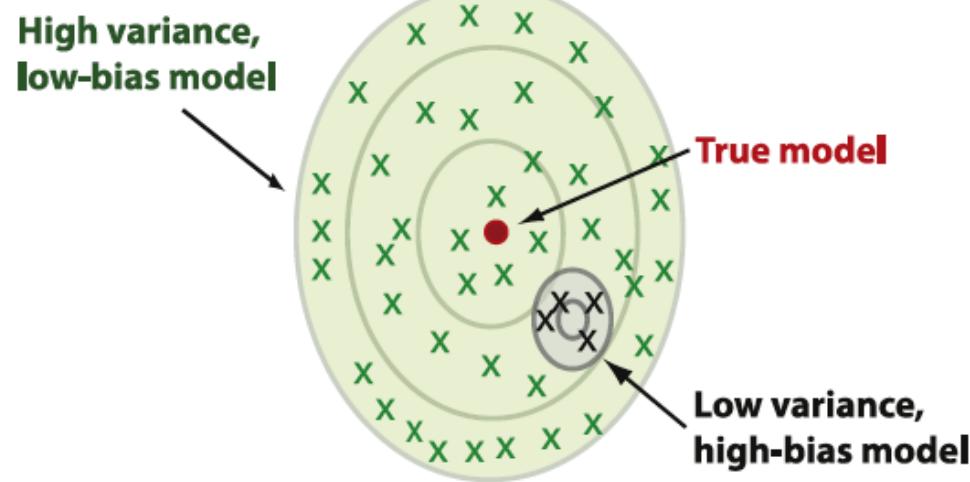
- Do you remember the Bayes error rate?
 - the lowest possible error rate given the features.
 - analogous to the irreducible error.

Bia-Variance Tradeoff

- **Variance** refers to the amount by which \hat{f} would **change** if we estimated it using a **different training data set**.
High variance -- small changes in the training data can result in large changes in \hat{f} .
- **Bias** refers to the error that is introduced by **approximating a real-life problem**, which may be **extremely complicated**, by a **much simpler model**.
The more we allow the model to accommodate to the training data set, the lower the bias.
- Example: the **flexible green curve**
(High Variance and Low Bias → Overfitting)

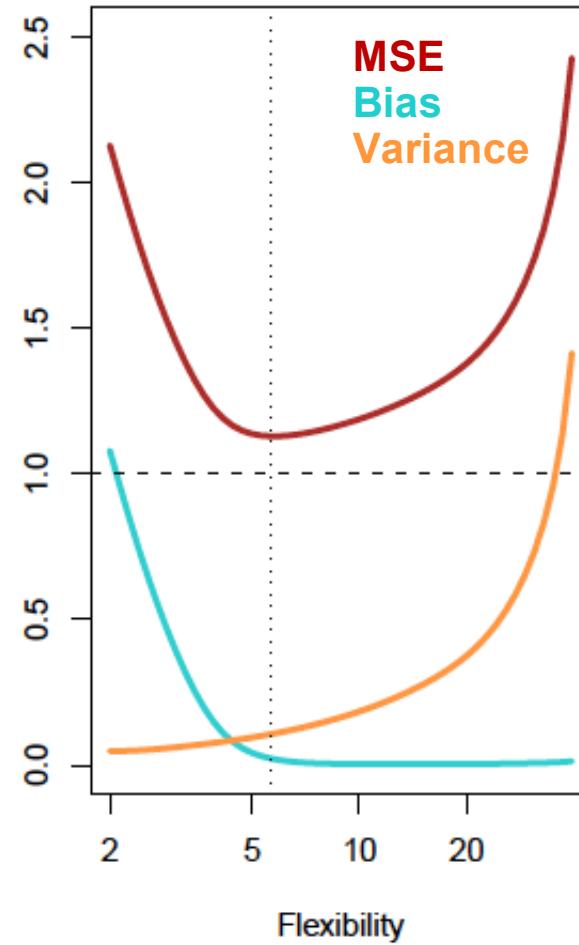
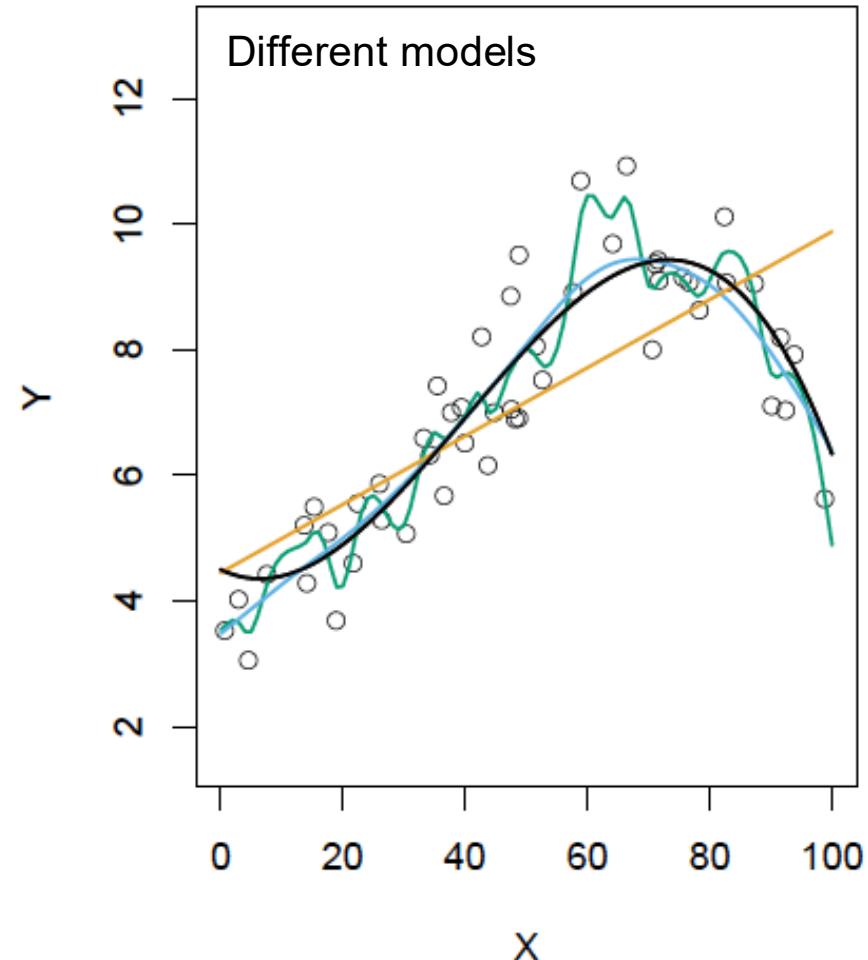


Bia-Variance Tradeoff



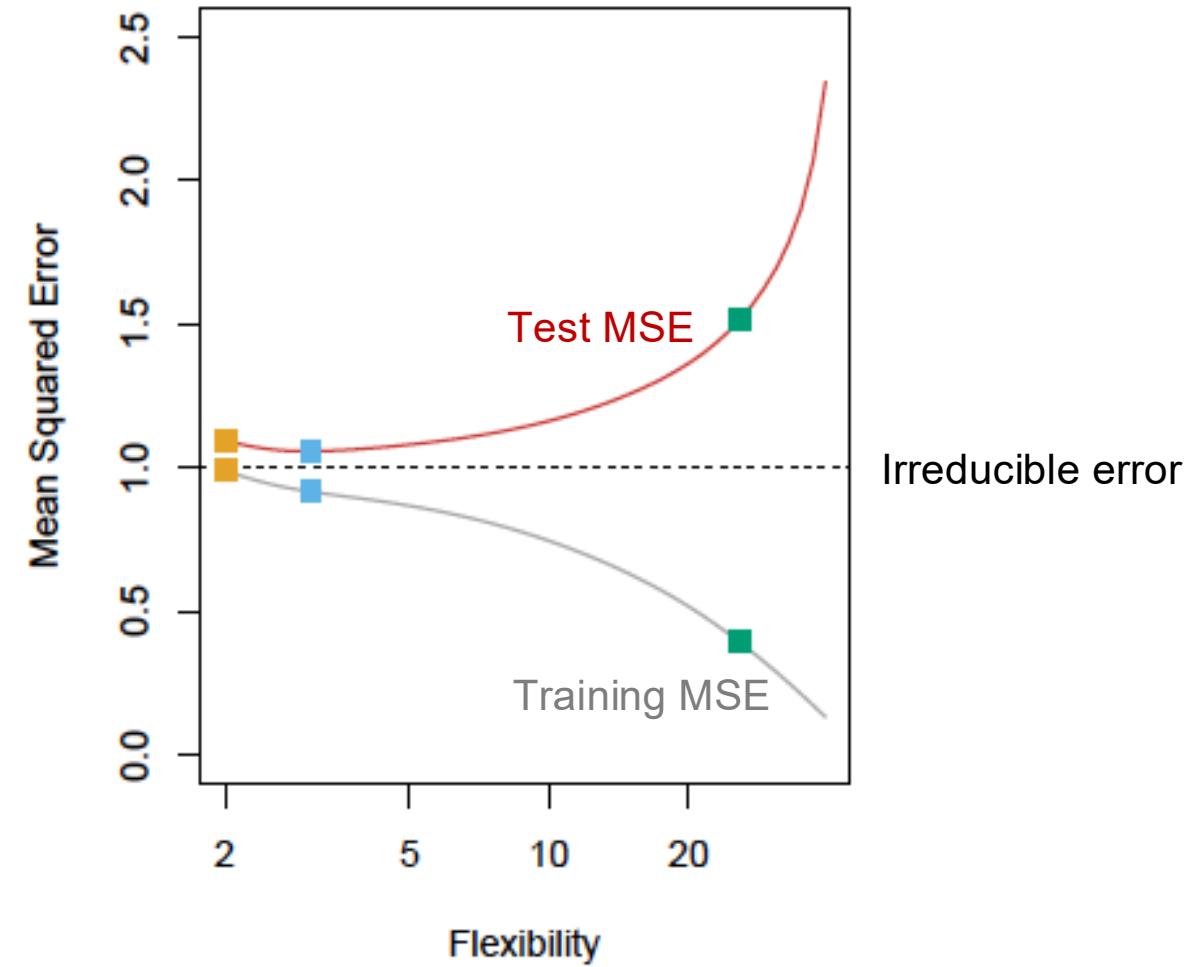
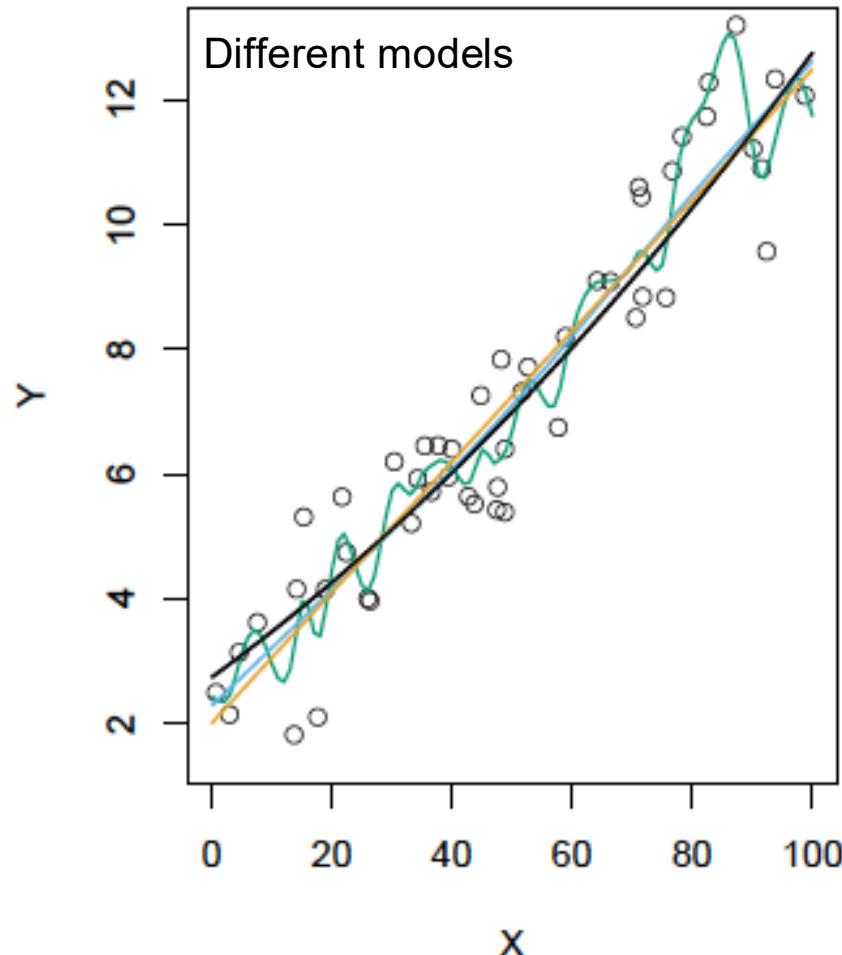
- Each point represents a model that has been fitted using some training data.

Bia-Variance Tradeoff

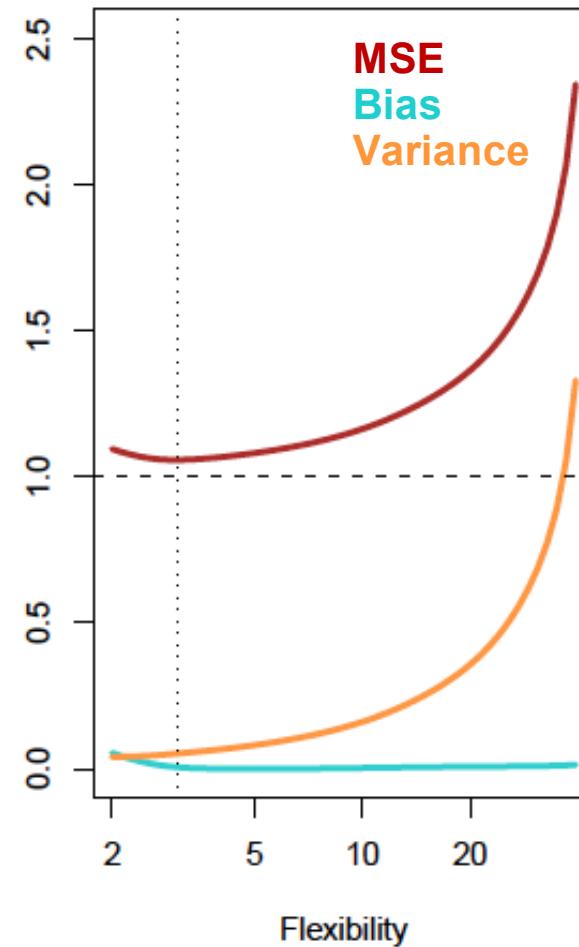
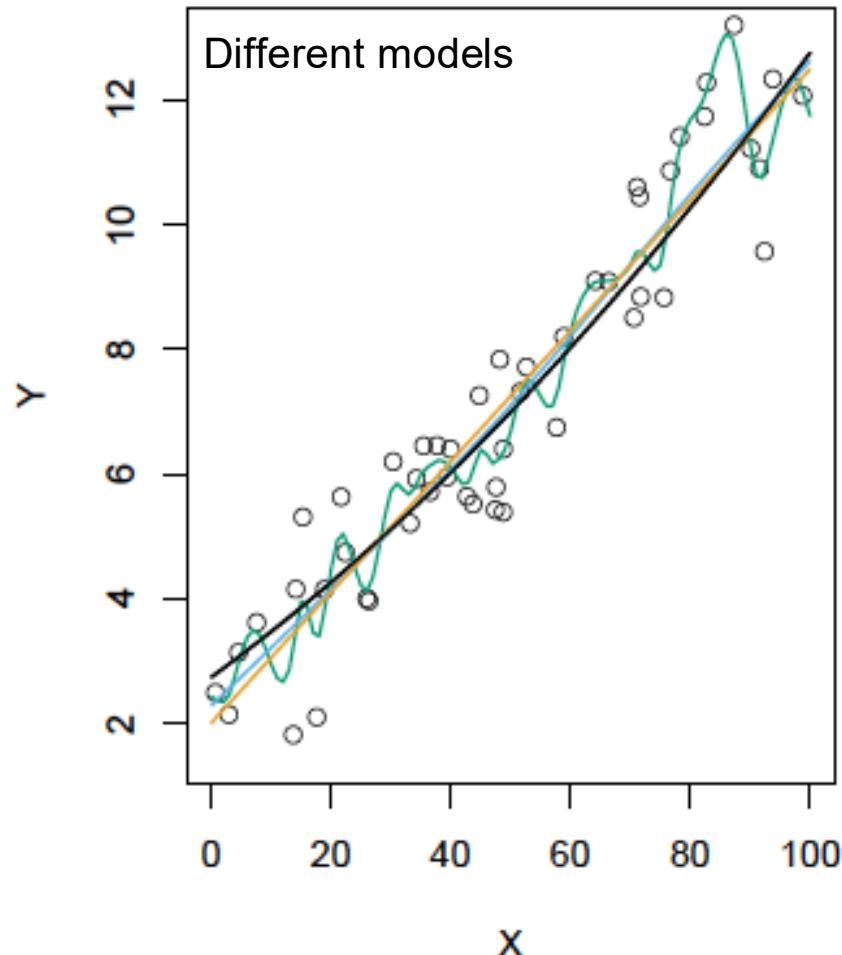


Irreducible error

Bia-Variance Tradeoff



Bia-Variance Tradeoff



Regularization

Ridge

- Avoid overfitting to the observed data, especially when working with a small data set.
- Worsen the predictions by punishing the model:

$$\begin{aligned}\hat{\beta}_{\text{Ridge}} &= \underset{\beta}{\operatorname{argmin}} \sum_{n=1}^N (y_n - x_n^T \beta)^2 + \lambda \beta^T \beta \quad \text{L2} \\ &= \underset{\beta}{\operatorname{argmin}} (y - X\beta)^T (y - X\beta) + \lambda \beta^T \beta\end{aligned}$$

- Ridge regression has a closed-form solution:

$$\hat{\beta}_{\text{Ridge}} = (X^T X + \lambda I)^{-1} X^T y$$

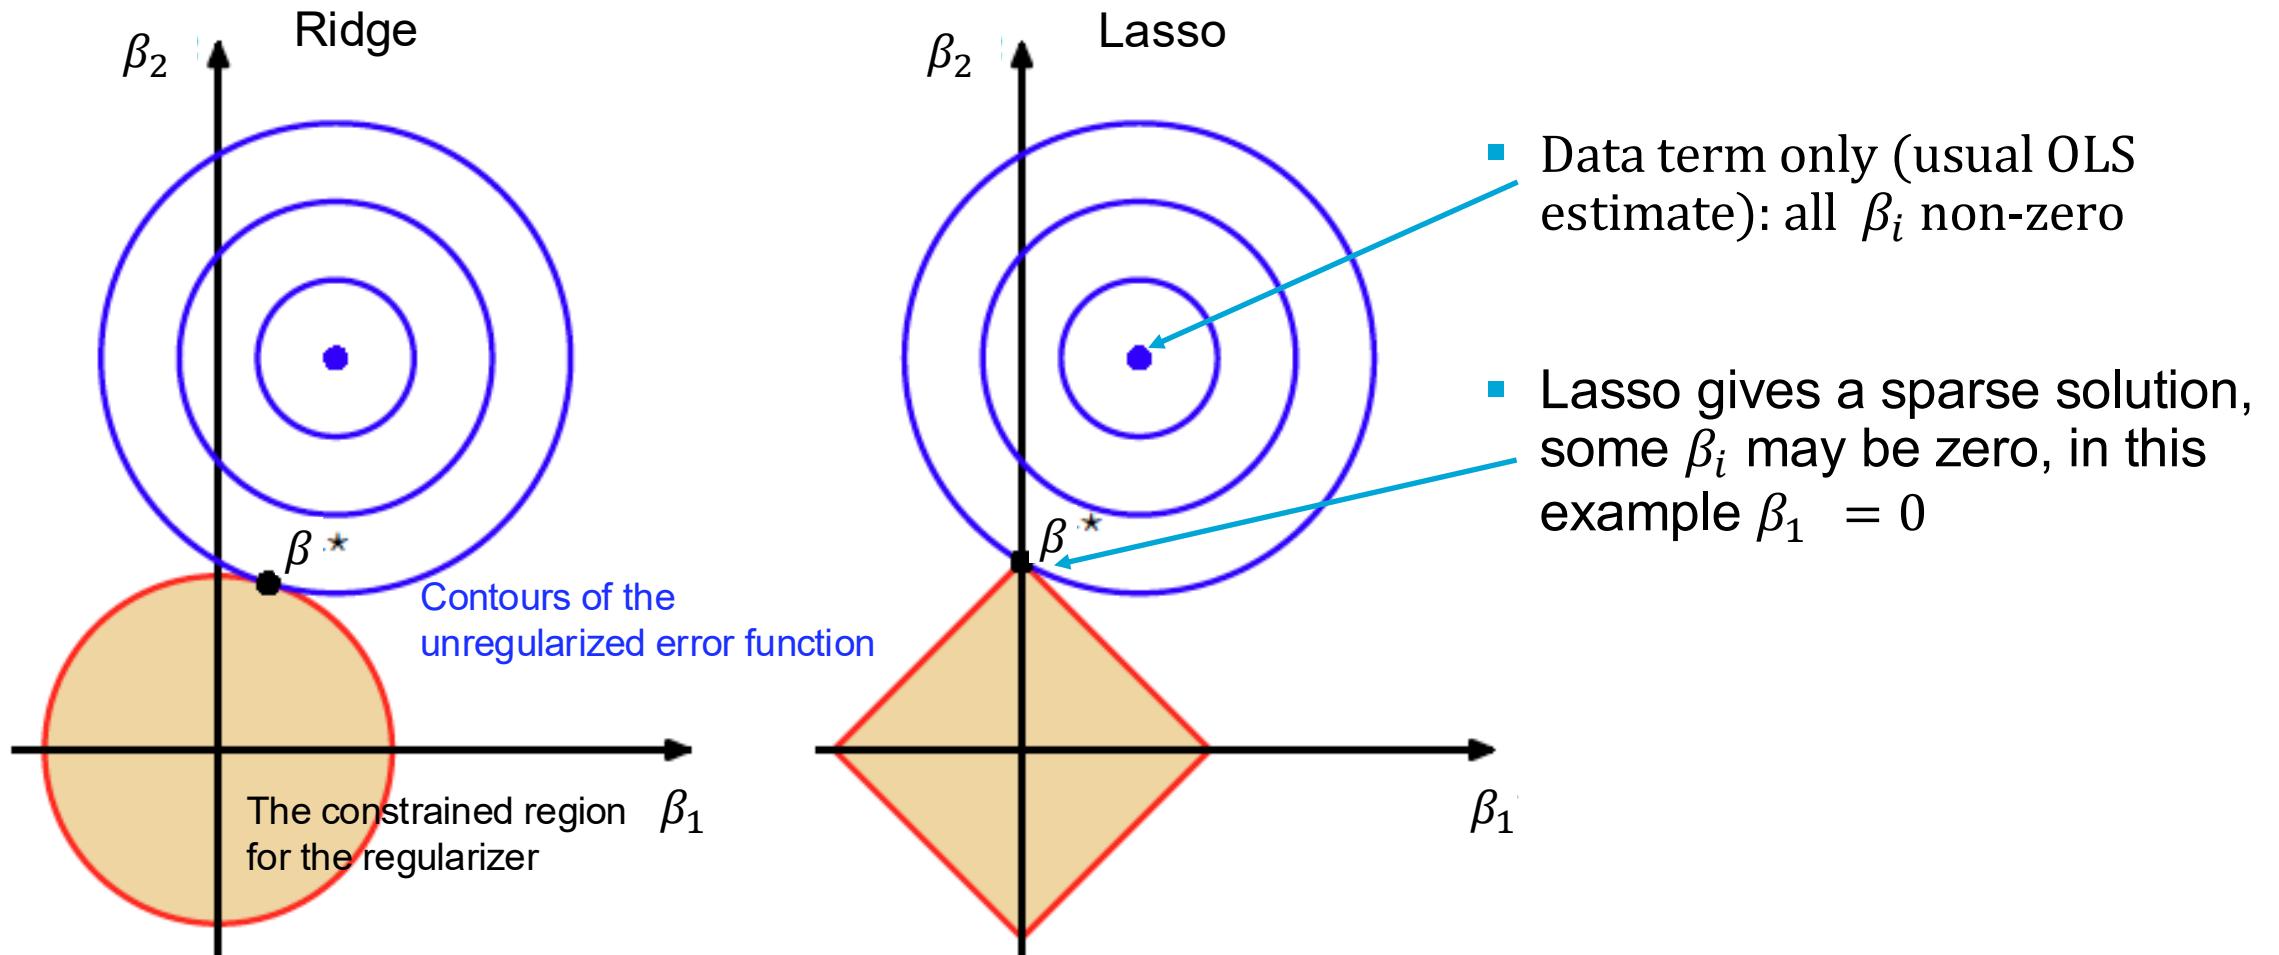
Lasso

- Lasso: least absolute shrinkage and selection operator

$$\begin{aligned}\hat{\beta}_{\text{Lasso}} &= \underset{\beta}{\operatorname{argmin}} \sum_{n=1}^N (y_n - x_n^T \beta)^2 + \lambda \|\beta\|_1 \\ &= \underset{\beta}{\operatorname{argmin}} (y - X\beta)^T (y - X\beta) + \lambda \|\beta\|_1\end{aligned}$$

- No closed-form solution.
(e.g., coordinate descent. “The elements of statistical learning” Chapter 3.4.4 & 3.8.6 if interested)
- Introduce sparsity to the parameter space.
- Give penalty for having many predictors with small effects.

Regularizations



Classifiers: Logistic Regression and Perceptron

Logistic Regression

- Logistic regression model arises from the desire to model the **posterior probabilities of the K classes** via linear functions in x , while at the same time ensuring that they sum to one and remain in $[0, 1]$.

- -- Two classes



- -- More than two classes

The Bayes Rule if you still remember...

$$p(A|B) = \frac{\text{likelihood} \quad \text{prior}}{\text{evidence}}$$
$$p(A|B) = \frac{p(B|A)p(A)}{p(B)}$$

Logistic Regression

- Consider the case of two classes,

$$p(C_1|x) = \frac{p(x|C_1)p(C_1)}{p(x|C_1)p(C_1) + p(x|C_2)p(C_2)}$$

- We define $a = \ln\left(\frac{p(C_1|x)}{p(C_2|x)}\right) = \ln\frac{p(x|C_1)p(C_1)}{p(x|C_2)p(C_2)}$

The inverse of the logistic sigmoid, i.e.,
the logit function, a.k.a. the log odds.

$$p(C_1|x) = \frac{1}{1 + \exp(-a)} = \sigma(a)$$

*The logistic **sigmoid** (S-shaped) function.*

Logistic Regression

- Consider the case $K > 2$ classes,

$$p(C_k|x) = \frac{p(x|C_k)p(C_k)}{\sum_j p(x|C_j)p(C_j)}$$

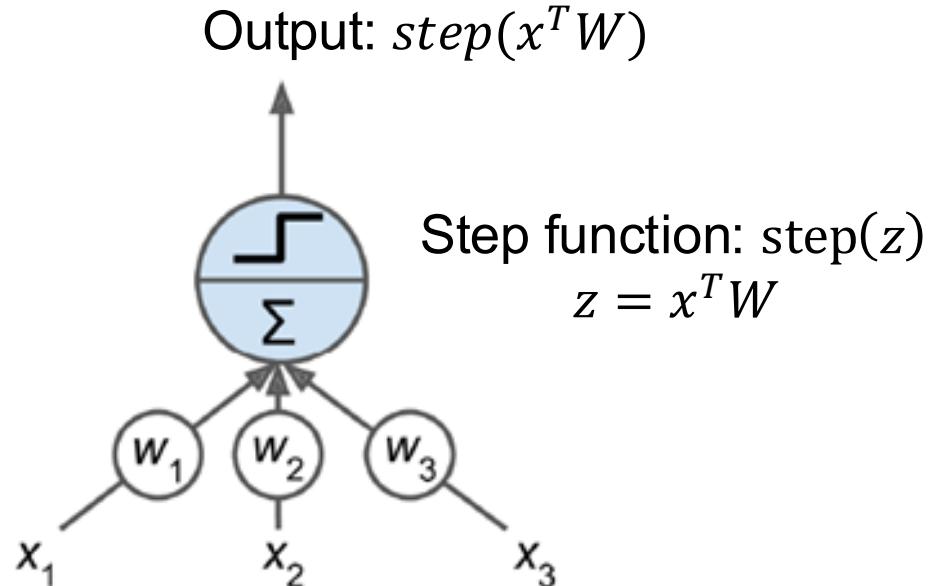
$$a_k = \ln p(x|C_k)p(C_k)$$

$$p(C_k|x) = \frac{\exp(a_k)}{\sum_j \exp(a_j)} = \sigma(a)$$

- The *normalized exponential*, a.k.a. *the softmax function*, is a multiclass generalization of *the logistic sigmoid*.

Perceptron

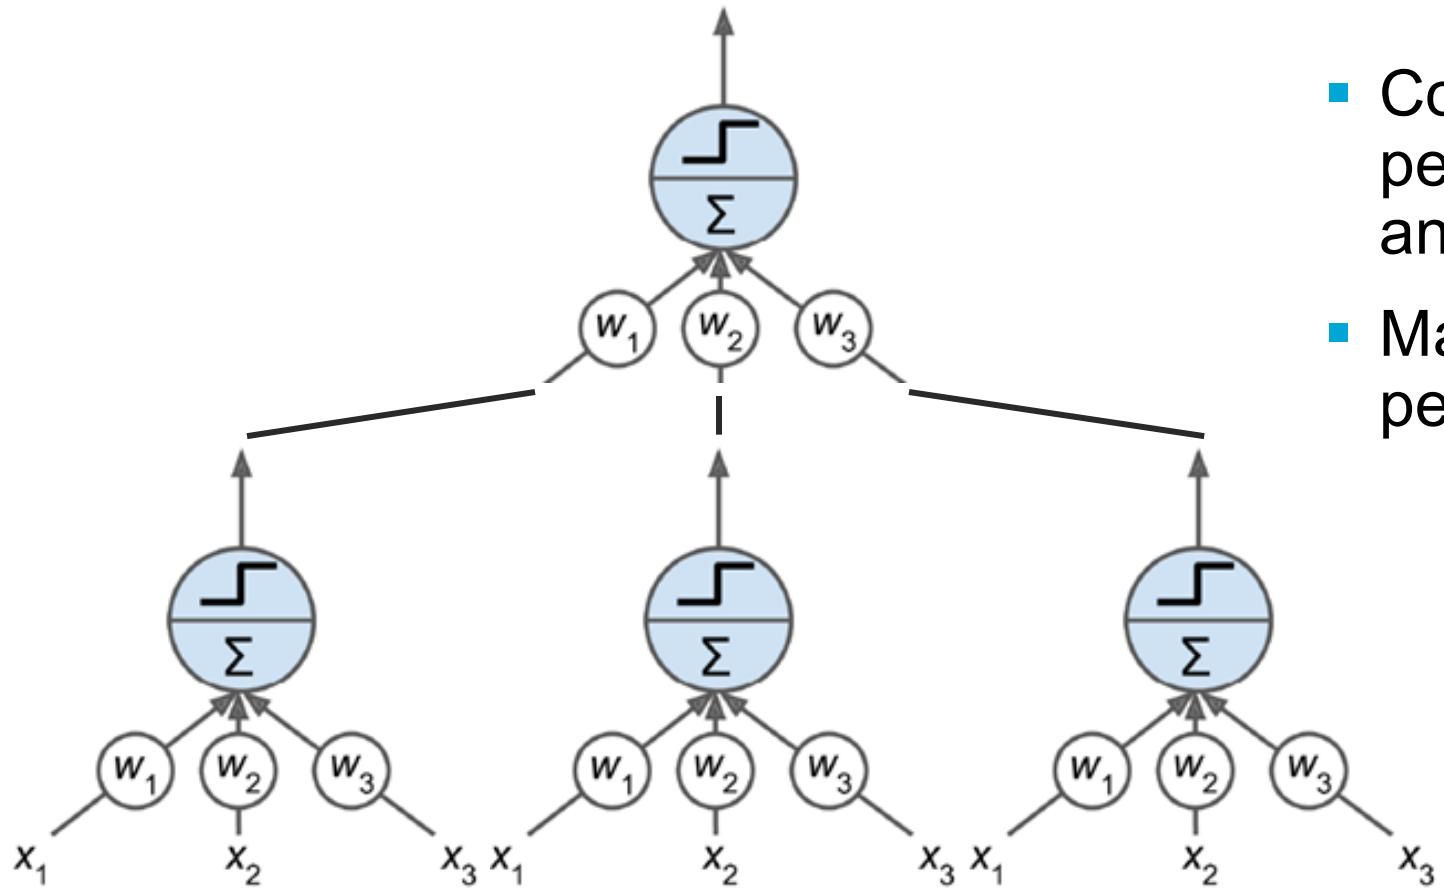
- Invented in 1957 by Frank Rosenblatt
- Linear classifier combined with a Heaviside/unit step function
- Step functions a.k.a. non-linearities



$$Heaviside(z) = \begin{cases} 0 & \text{if } z < 0 \\ 1 & \text{if } z \geq 0 \end{cases}$$

- For regression?
- What if we use the sigmoid function as the step function?

Multilayer Perceptron (MLP)



- Connect the output of one perceptron to the inputs of another perceptron
- Make multiple layers, stacking perceptrons on top of each other.

➤ For multiclass classification?
Softmax?

Bayesian Linear Regression

Look at OLS again from a probabilistic perspective

- We assume:

$$y_n = x_n^T \beta + \epsilon, \quad \text{where } \epsilon \text{ i.i.d with } \mathcal{N}(0, \sigma_\epsilon^2)$$

- Making assumption that data points are drawn independently from the above distribution, the likelihood function is:

$$\begin{aligned} p(y|X, \beta, \sigma_\epsilon^2) &= \prod_{n=1}^N p(y_n|x_n^T \beta, \sigma_\epsilon^2) \\ &= \frac{1}{\sqrt{(2\pi)^N \sigma_\epsilon^{2N}}} \exp\left(-\frac{1}{2\sigma_\epsilon^2} (y - X\beta)^T (y - X\beta)\right) \end{aligned}$$

Look at OLS again from a probabilistic perspective

- Estimate parameters where $p(y|X, \beta, \sigma_\epsilon^2)$ is maximized:

$$\hat{\beta}_{ML} = \operatorname{argmax}_\beta p(y|X, \beta, \sigma_\epsilon^2)$$

- We get:

$$\hat{\beta}_{ML} = (X^T X)^{-1} X^T y$$

- Similarly, we can maximize the log likelihood function with respect to the σ_ϵ^2 .

$$\hat{\sigma}_{ML}^2 = \frac{1}{N} \sum_{n=1}^N (r_n)^2$$

Look at OLS again from a probabilistic perspective

- Recall

$$\hat{\beta}_{OLS} = (X^T X)^{-1} X^T y$$

- We see

$$\hat{\beta}_{ML} = \hat{\beta}_{OLS}$$

- OLS could be motivated as the **maximum likelihood** solution under an assumed Gaussian noise.
- What is the limitation? -- There is no representation of our **uncertainty**.

Bayesian Linear Regression

- The aim is not to find the best “single” value for the parameters, but rather to determine their posterior distribution.
- Recall the Bayes rule again,

$$p(A|B) = \frac{\underset{likelihood}{p(B|A)} \underset{prior}{p(A)}}{\underset{evidence}{p(B)}}$$

$$p(\beta|D) = \frac{p(D|\beta)p(\beta)}{p(D)} \quad D = ((x_1, y_1), (x_2, y_2), \dots, (x_N, y_N)), x_n \in \mathbb{R}^d, y_n \in \mathbb{R}$$

Maximum a Posterior (MAP) Estimation

- We have $y \sim \mathcal{N}(X\beta, \sigma_\epsilon^2 I)$

$$\hat{\beta}_{MAP} = \operatorname{argmax}_\beta p(\beta | X, y)$$

Likelihood of the response given the predictors and the model

Prior probability of the model parameters (we can use prior such as Gaussian)

$$= \operatorname{argmax}_\beta \frac{p(y|X, \beta) p(\beta)}{p(y|X)}$$

a normalization constant

$$= \operatorname{argmax}_\beta p(y|X, \beta)p(\beta)$$

$$\hat{\beta}_{MAP} = \left(\frac{\sigma_\epsilon^2}{\sigma_\beta^2} I + X^T X \right)^{-1} X^T y$$

Try the derivation!