




Dimensionality Reduction Feature Selection & Extraction

Jing Sun

Is “Income” an informative feature?




Feature Space


- A p -dimensional space, in which each dimension is a feature containing n [labeled] samples [objects]

- What will happen if p is very large?
- **[the curse of dimensionality]**

- In high-dimensional spaces, our 2D/3D intuition does not work anymore...



High-Dimensional Spaces




- Example:
- Neighborhood capturing 10% of uniformly distributed data in hypercube
- E.g. in \mathbb{R}^{20} side length of $\sqrt[20]{1} \approx 0.89$
 - So, not a small block anymore...

High-Dimensional Spaces

- Example: Boundary points ?

500 samples from normal distribution

In a 2-D space,
only 2% are on the convex hull




High-Dimensional Spaces

- Example: Boundary points ?

500 samples from normal distribution

In a 2-D space,
only 2% are on the convex hull

In a 20-D space,
95% are on the convex hull !



High-Dimensional Spaces


- Example: Points tend to have equal distances

200 samples from normal distribution

$$N(2000, 8000)$$

In a \mathbb{R}^1 to \mathbb{R}^{1000} space

Consider $\frac{\text{std}(d^2)}{\text{mean}(d^2)}$ for squared distance d^2




High-Dimensional Spaces

- Example: Points tend to have equal distances

200 samples from normal distribution

$$N(2000, 8000)$$


In a \mathbb{R}^1 to \mathbb{R}^{1000} space



Dimensionality Reduction

- Problem: **too few samples in too many dimensions**
[the curse of dimensionality]

- Solution: drop dimensions / features
 - Feature selection
 - Feature extraction




Dimensionality Reduction

- Uses/Benefits :
 - Fewer parameters give **faster** algorithms and parameters are **easier** to estimate
 - **Explaining** which measurements are useful and which are not [**reducing redundancy**]
 - **Visualization of data** can be a powerful tool when designing pattern recognition systems

Dimensionality Reduction

- Problem: **too few samples in too many dimensions**
[the curse of dimensionality]
- Solution: drop dimensions / features
 - Feature selection
 - Feature extraction
- **Questions:**
 - Which dimensions to drop?
 - What feature subset to keep?





Dimensionality Reduction by Selection or Extraction

- Overview – **Feature Selection** vs **Feature Extraction**
- Criteria
 - Mahalanobis distance (vs Euclidean distance)
 - Scatter matrices (what are S_W , S_B , S_T ?)
- Approaches
 - Sequential **feature selection** (individual, forward, backward, etc.)
 - Principal Component Analysis & Recall LDA (\in linear **feature extraction**)

Dimensionality Reduction by Selection or Extraction

- Overview – **Feature Selection vs Feature Extraction**
- Criteria
 - Mahalanobis distance (vs Euclidean distance)
 - Scatter matrices (what are S_W , S_B , S_T ?)
- Approaches
 - Sequential **feature selection** (individual, forward, backward, etc.)
 - Principal Component Analysis & Recall LDA (\in linear **feature extraction**)

Feature Selection vs Extraction




Feature Selection vs Extraction

- Feature selection :

Select d out of p measurements

Only a subset of the original features are selected.


There are $\binom{p}{d} = \frac{p!}{d!(p-d)!}$ subsets.



-
- Feature extraction :

Map p measurements to d measurements

All original features are used (they are transferred)




Dimensionality Reduction by Selection or Extraction

- Overview – **Feature Selection** vs **Feature Extraction**
- Criteria
 - Mahalanobis distance (vs Euclidean distance)
 - Scatter matrices (what are S_W , S_B , S_T ?)
- Approaches
 - Sequential **feature selection** (individual, forward, backward, etc.)
 - Principal Component Analysis & Recall LDA (\in linear **feature extraction**)

Why Mahalanobis distance?


- When measuring the distance from a **single point** to another **single point**, using (squared) Euclidean distance is fine.

$$D_E = (x_{red} - x_{yellow})^2 + (y_{red} - y_{yellow})^2$$




Why Mahalanobis distance?

- However,
- when there is a group of data points:
- **Centroid (mean vector)** = $\begin{pmatrix} \bar{x} \\ \bar{y} \end{pmatrix}$
- Euclidean distances $D_{E1} = D_{E2}$



Mahalanobis distance


- Takes the **variance** into account.
 - It is a distance measure between **a point** and **a distribution**.
 - For red and blue points,
- $$D_M = \begin{pmatrix} x - \bar{x} \\ y - \bar{y} \end{pmatrix}^T \Sigma^{-1} \begin{pmatrix} x - \bar{x} \\ y - \bar{y} \end{pmatrix}$$
- You will see $D_{M2} > D_{M1}$



Mahalanobis distance

- Think about:
- What if Σ is an identity matrix?

$$D_M = \begin{pmatrix} x - \bar{x} \\ y - \bar{y} \end{pmatrix}^T \mathbf{I} \begin{pmatrix} x - \bar{x} \\ y - \bar{y} \end{pmatrix} = D_E$$



Mahalanobis distance

- Mahalanobis distance between two classes:
 - Assumes Gaussian distributions with equal covariance matrix
 - $D_M = (\boldsymbol{\mu}_1 - \boldsymbol{\mu}_2)^T \mathbf{S}_W^{-1} (\boldsymbol{\mu}_1 - \boldsymbol{\mu}_2)$
- E.g., Exercise 6.21
- What is this \mathbf{S}_W ?

Dimensionality Reduction by Selection or Extraction

- Overview – **Feature Selection** vs **Feature Extraction**
- **Criteria**
 - Mahalanobis distance (vs Euclidean distance)
 - Scatter matrices (what are S_W , S_B , S_T ?)
- Approaches
 - Sequential **feature selection** (individual, forward, backward, etc.)
 - Principal Component Analysis & Recall LDA (\in linear **feature extraction**)

Scatter Matrices

- Within-class scatter matrix:

$$S_W = \sum_{i=1}^M \frac{n_i}{N} \Sigma_i, \quad \Sigma_i \text{ is the covariance matrix of class } w_i; M \text{ is the number of classes;} \\ n_i \text{ is the number of samples in class } w_i, \text{ out of a total of } N \text{ samples.}$$

- Between-class scatter matrix:

$$S_B = \sum_{i=1}^M \frac{n_i}{N} (\mu_i - \mu)(\mu_i - \mu)^T, \quad \mu_i \text{ is the mean of class } w_i, \\ \mu \text{ is the global mean.}$$




$$\mu = \sum_{i=1}^M \frac{n_i}{N} \mu_i$$

- Total scatter matrix: $S_T = S_W + S_B$

Scatter Matrices

For a classification task




- S_W = “average class width”; ***the smaller, the better***
- S_B = “average distance between class means”; ***the larger, the better***
- S_T = “overall width”



Scatter Matrices

For a classification task

- S_W = “average class width”; ***the smaller, the better***
- S_B = “average distance between class means”; ***the larger, the better***
- S_T = “overall width”




Scatter-based Criteria

- $J_1 = \frac{\text{trace}\{S_T\}}{\text{trace}\{S_W\}}$
- $J_2 = \frac{|S_T|}{|S_W|}$
- etc.
- by using **various combinations** of S_W , S_B , S_T in a “trace” or “determinant” formulation...
- PS: The “trace” is equal to the sum of the eigenvalues; the “determinant” is equal to their product.

FDR: Fisher Discriminant Ratio

- 1-D, two-class problem
- $S_W \propto (\sigma_1^2 + \sigma_2^2)$, $S_B \propto (\mu_1 - \mu_2)^2$,
- Combining S_W and S_B , you get Fisher's criterion

$$J_F = \frac{(\mu_1 - \mu_2)^2}{\sigma_1^2 + \sigma_2^2}$$



- It is often used to quantify the separability capabilities of individual features.

Dimensionality Reduction by Selection or Extraction

- Overview – **Feature Selection** vs **Feature Extraction**
- Criteria
 - Mahalanobis distance (vs Euclidean distance)
 - Scatter matrices (what are S_W , S_B , S_T ?)
- Approaches
 - Sequential **feature selection** (individual, forward, backward, etc.)
 - Principal Component Analysis & Recall LDA (\in linear **feature extraction**)

Which method would guarantee optimal performance?

- I have p features (let's say $p = 40$).
- I think this is too many to handle...
- I want to select d features from p .
- But what should $d = ?$ Well, I'm not sure...
- What can I do?

Which method would guarantee optimal performance?

- Trying all possible feature combinations



- **Exhaustive** feature selection

$$\binom{p}{d} = \frac{p!}{d!(p-d)!}$$

$$\sum_{i=1}^p \binom{p}{i} \text{ combinations}$$

- If originally there are 4 features, we will end up with 15 combinations.

- $\binom{4}{1} + \binom{4}{2} + \binom{4}{3} + \binom{4}{4} = 15$

- But, what if there are 40 features...?

-- over a billion! :-)

Sub-optimal Strategies


- Trying all possible feature combinations
- **Exhaustive** feature selection
- It can be super Expensive! And Exhaustive!! :-(
- Let's use **Sequential Feature Selection!**



$$\binom{p}{d} = \frac{p!}{d!(p-d)!}$$


$$\sum_{i=1}^p \binom{p}{i}$$
 combinations

Feature Selection Methods




Forward Selection (FS)

- Start with **empty feature set**



Forward Selection (FS)


- Start with empty feature set



- Compute the **criterion value** for each feature individually and **select the best one**,

Forward Selection (FS)


- Start with empty feature set



- Compute the **criterion value** for each feature individually and **select the best one**,
 $x_2 > x_4 > x_2 > x_3 \Rightarrow x_2$

Forward Selection (FS)


- Start with empty feature set



- Compute the criterion value for each feature individually and select the best one,
 $x_2 > x_4 > x_2 > x_3 \Rightarrow x_2$
- **Keep the winner** and compute the criterion value for **all two-feature combinations** that include it.

Forward Selection (FS)


- Start with empty feature set



- Compute the criterion value for each feature individually and select the best one,
 $x_2 > x_4 > x_2 > x_3 \Rightarrow x_2$
- **Keep the winner** and compute the criterion value for **all two-feature combinations** that include it.
 $[x_2, x_1] > [x_2, x_4] > [x_2, x_3]$

Forward Selection (FS)


- Start with empty feature set



- Compute the criterion value for each feature individually and select the best one,
 $x_2 > x_4 > x_2 > x_3 \Rightarrow x_2$
 - Keep the winner and compute the criterion value for all two-feature combinations that include it.
 $[x_2, x_1] > [x_2, x_4] > [x_2, x_3]$
- ... until a **predefined number** of features are left.


Backward Selection (BS)

- Start with **all originally available features**



Backward Selection (BS)

- Start with all originally available features



- Compute the criterion value for **all possible combinations after eliminating one feature**,

Backward Selection (BS)


- Start with all originally available features



- Compute the criterion value for **all possible combinations after eliminating one feature**,
 $[x_1, x_2, x_4] > [x_1, x_2, x_3] > [x_2, x_3, x_4] > [x_1, x_3, x_4]$
Keep the winner combination (i.e., remove one feature);

Backward Selection (BS)

- Start with all originally available features



- Compute the criterion value for all possible combinations after eliminating one feature,
 $[x_1, x_2, x_4] > [x_1, x_2, x_3] > [x_2, x_3, x_4] > [x_1, x_3, x_4]$
Keep the winner combination (i.e., remove one feature);
- Repeat step above: **from the winner vector, eliminate one feature**, and for **each of the resulting combinations**, compute the criterion value... ...

Backward Selection (BS)

- Start with all originally available features




- Compute the criterion value for all possible combinations after eliminating one feature,
 $[x_1, x_2, x_4] > [x_1, x_2, x_3] > [x_2, x_3, x_4] > [x_1, x_3, x_4]$
Keep the winner combination (i.e., remove one feature);
- Repeat step above: **from the winner vector, eliminate one feature**, and for **each of the resulting combinations**, compute the criterion value... ...
 $[x_1, x_2] > [x_2, x_4] > [x_1, x_4]$
... until a **predefined number** of features are left.

Bidirectional Selection

- It **applies FS and BS simultaneously**:
 - FS starts from the empty feature set.
 - BS starts from the full set of all originally available features.
- To make sure they **converge to the same solution**:
 - Features already selected by FS are **not removed** by BS.
 - Features already removed by BS are **not selected** by FS.


Bidirectional Selection

Four features in order of x_1, x_2, x_3, x_4 ,
1 means selected, 0 means not selected,
e.g., (0,0,0,1) means only x_4 is selected.



Bidirectional Selection

Four features in order of x_1, x_2, x_3, x_4 ,
1 means selected, 0 means not selected,
e.g., (0,0,0,1) means only x_4 is selected.



BS Full set of all originally available features


x_1, x_2, x_4

x_2

FS Empty feature set

Bidirectional Selection

Four features in order of x_1, x_2, x_3, x_4 ,
1 means selected, 0 means not selected,
e.g., (0,0,0,1) means only x_4 is selected.



BS Full set of all originally available features

x_1, x_2, x_4

x_2, x_4

x_2

FS Empty feature set

Plus- L Take-away- R Selection

- **Also based on the ideas of FS and BS.** It has two forms.
- If $L > R$, it starts from the **empty set** and
 - repeatedly add L features
 - repeatedly remove R features
- If $L < R$, it starts from the **full set** of all available features and
 - repeatedly remove R features
 - repeatedly add L features
- There is no way of foreseeing the best values of L and R . :-(

Floating Selection

- FS and BS suffer from the so-called **nesting effect**. That is,
 - For FS, once a feature is chosen, there is **no way** for it to be discarded later on.
 - For BS, once a feature is discarded, there is **no way** for it to be reconsidered again.
- Plus- L Take-away- R Selection doesn't have a flexible backtracking capability.
 - Every round, we **have to** plus L and **have to** take away R .
- **Floating Selection** allows flexible backtracking:
 - The **dimensionality of the subset** during the search can be "**floating**" up and down.

Floating Selection

- There are two floating methods:
 - Floating **forward** selection & Floating **backward** selection
- Floating **forward** selection starts from the **empty set**,
 - after each forward step, it performs backward steps as long as the criterion function increases.
- Floating **backward** selection starts from the **full set**,
 - after each backward step, it performs forward steps as long as the criterion function increases.

Dimensionality Reduction by Selection or Extraction

- Overview – **Feature Selection** vs **Feature Extraction**
- Criteria
 - Mahalanobis distance (vs Euclidean distance)
 - Scatter matrices (what are S_W , S_B , S_T ?)
- Approaches
 - Sequential **feature selection** (individual, forward, backward, etc.)
 - Principal Component Analysis & Recall LDA (\in linear **feature extraction**)

Interesting facts about PCA

- PCA is widely recognized as the most classical method for dimensionality reduction, having been invented in 1901.
- However, it **doesn't automatically reduce the dimensionality!**
- Rather, it **transforms the data into a new coordinate system** where **the choice to retain fewer** principal components effectively reduces dimensionality.
 - **Retain the variance as much as possible**
 - i.e., **Minimize the reconstruction error**

PCA: offers different view of your data

- Data:

$$\mathbf{X} = \begin{bmatrix} \mathbf{x}_1 \\ \mathbf{x}_2 \\ \vdots \\ \mathbf{x}_p \end{bmatrix}$$

mean-centered data (the mean of each feature is 0);
 p is number of features

- (Variance-) Covariance matrix:

$$\Sigma = \mathbf{X}\mathbf{X}^T = \begin{bmatrix} \sigma_1^2 & \sigma_{12} & \dots & \sigma_{1p} \\ \sigma_{21} & \sigma_2^2 & \cdots & \sigma_{2p} \\ \vdots & \vdots & \ddots & \vdots \\ \sigma_{p1} & \sigma_{p2} & \dots & \sigma_p^2 \end{bmatrix}$$

$(p \times p)$

PCA: offers different view of your data

- Eigen-decomposition of the covariance matrix:

$$\Sigma v = v\lambda, \|v\|^2 = 1 \longrightarrow v_i = \begin{matrix} v_{1i} \\ v_{2i} \\ \vdots \\ v_{pi} \end{matrix}, \lambda_i, i = 1, 2, \dots p$$

- Transform the data to a new space, in which the coordinate system is defined by the principal components.

$$V = [v_1 \ v_2 \ \dots \ v_k \ \dots \ v_p]_{(p \times p)}$$


Each column of V is a principal component
ORDERED by the value of λ ,
 λ_1 is the largest eigenvalue

$$T = V^T X = \begin{matrix} v_1^T X \\ v_2^T X \\ \vdots \\ v_k^T X \\ \vdots \\ v_p^T X \end{matrix}$$


Quiz:
What is the dimensionality of T ?
Same with X ?

PCA: offers different view of your data

Original Space (2D)



PCA Space (2D)




PCA: choose to reduce dimensionality

- Again, PCA doesn't automatically reduce the dimensionality.

$$\mathbf{T} = \mathbf{V}^T \mathbf{X}$$

- Choose to retain** the first k principal components because e.g., 95% variance is captured



$$\mathbf{V} = [\mathbf{v}_1 \ \mathbf{v}_2 \dots \mathbf{v}_k \ \boxed{\dots \mathbf{v}_p}]$$


What is the dimensionality of \mathbf{T}_k ?

$$\mathbf{T}_k = \mathbf{V}_k^T \mathbf{X}$$

$$\mathbf{V}_k^T = \begin{matrix} (k \times p) \\ \left[\begin{matrix} \mathbf{v}_1^T \\ \mathbf{v}_2^T \\ \vdots \\ \mathbf{v}_k^T \end{matrix} \right] \end{matrix}$$

PCA: choose to reduce dimensionality



Quiz

$$V = \begin{bmatrix} v_1 & v_2 & \dots & v_k & \dots & v_p \end{bmatrix}_{(p \times p)}$$

$$T_k = V_k^T X = \begin{bmatrix} v_1^T X \\ v_2^T X \\ \vdots \\ v_k^T X \end{bmatrix}_{(k \leq p)}$$

- When $k = p$, T_k contain **exactly the same amount** of information as the original data X .
True or False?
- What does $v_1^T X$ in T_k represent?
- What does $v_1^T \Sigma v_1$ represent?

Quiz


$$V = [\boldsymbol{v}_1 \ \boldsymbol{v}_2 \dots \boldsymbol{v}_k \dots \boldsymbol{v}_p]_{(p \times p)}$$

$$\mathbf{T}_k = V_k^T X = \begin{bmatrix} \boldsymbol{v}_1^T X \\ \boldsymbol{v}_2^T X \\ \vdots \\ \boldsymbol{v}_k^T X \end{bmatrix}_{(k \leq p)}$$

- When $k = p$, \mathbf{T}_k contain **exactly the same amount** of information as the original data X .
True!!!
- What does $\boldsymbol{v}_1^T X$ in \mathbf{T}_k represent?
Projection onto the first principal component!
- What does $\boldsymbol{v}_1^T \Sigma \boldsymbol{v}_1$ represent?
Variance!


PCA on MNIST data

- PCA reconstructions
- Original space (784 D)



PCA on MNIST data


- 50% Variance: Dim = 11



- The more PCs we retain, the smaller the reconstruction error becomes.

PCA on MNIST data


- 95% Variance: Dim = 154



- The more PCs we retain, the smaller the reconstruction error becomes.

PCA on MNIST data

- 99.5% Variance: Dim = 331




- The more PCs we retain, the smaller the reconstruction error becomes.

PCA on MNIST data

784 PCs

Original




2 PCs

Compressed



PCA SPACE



Colours indicate the class of the object

Two classical linear feature extractors

- LDA (or Fisher mapping):

Supervised


-- Find projection vector a that captures the greatest **separability between the classes**, i.e., choose a to maximize Fisher criterion:

$$J_F(a) = \frac{a^T S_B a}{a^T S_W a}$$

- PCA:



Unsupervised

-- Capture the greatest **variance in the total data**



Kernel PCA – Non-linear Dimensionality Reduction.


- Linear subspaces may be inefficient for some cases.
- The data are not linearly separable in the original dimension.




Kernel PCA – What it does

- Use a kernel function to project data into a higher-dim. space where they are linearly separable.
- $\phi : R^2 \rightarrow R^3$ $(x_1, x_2) \longrightarrow (z_1, z_2, z_3) = (x_1^2, \sqrt{2}x_1x_2, x_2^2)$

Data in 2D space




Data mapped to 3D space




Kernel PCA – What it does

Apply PCA



Apply Kernel PCA and drop PC3



Practice

Given mean-centered data in 3D for which the covariance matrix is given by

$$\Sigma = \begin{pmatrix} 1 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 4 \end{pmatrix}.$$

Also given is a data transformation matrix

$$R = \begin{pmatrix} 1 & 0 & 0 \\ 0 & \frac{1}{2} & -\frac{\sqrt{3}}{2} \\ 0 & \frac{\sqrt{3}}{2} & \frac{1}{2} \end{pmatrix},$$

by which we can linearly transform every data vector x (taken as a column vector) to a new 3D column vector z through $z = Rx$.

We note that R is actually a rotation matrix that rotates in the second and third coordinate.

Also note that for its inverse, we have

$$R^{-1} = \begin{pmatrix} 1 & 0 & 0 \\ 0 & \frac{1}{2} & \frac{\sqrt{3}}{2} \\ 0 & -\frac{\sqrt{3}}{2} & \frac{1}{2} \end{pmatrix}.$$

Q1: What is the first principal component of the original data for which we have the covariance matrix Σ ?

Q2: Assume we transform all the data by the transformation matrix R , what does the covariance of the transformed data become?

Q3: What is the first principal component for the transformed data?

Practice

Given mean-centered data in 3D for which the covariance matrix is given by

$$\Sigma = \begin{pmatrix} 1 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 4 \end{pmatrix}.$$

Also given is a data transformation matrix

$$R = \begin{pmatrix} 1 & 0 & 0 \\ 0 & \frac{1}{2} & -\frac{\sqrt{3}}{2} \\ 0 & \frac{\sqrt{3}}{2} & \frac{1}{2} \end{pmatrix},$$

by which we can linearly transform every data vector x (taken as a column vector) to a new 3D column vector z through $z = Rx$.

$$\text{Q2: } R\Sigma R^T = \begin{bmatrix} 1 & 0 & 0 \\ 0 & \frac{7}{2} & \frac{-\sqrt{3}}{2} \\ 0 & \frac{-\sqrt{3}}{2} & \frac{5}{2} \end{bmatrix}$$

We note that R is actually a rotation matrix that rotates in the second and third coordinate.

Also note that for its inverse, we have

$$R^{-1} = \begin{pmatrix} 1 & 0 & 0 \\ 0 & \frac{1}{2} & \frac{\sqrt{3}}{2} \\ 0 & -\frac{\sqrt{3}}{2} & \frac{1}{2} \end{pmatrix}.$$

$$\boldsymbol{v}_1 = \begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix}$$

Q1: What is the first principal component of the original data for which we have the covariance matrix Σ ?

Q2: Assume we transform all the data by the transformation matrix R , what does the covariance of the transformed data become?

Q3: What is the first principal component for the transformed data?

$$R\boldsymbol{v}_1$$